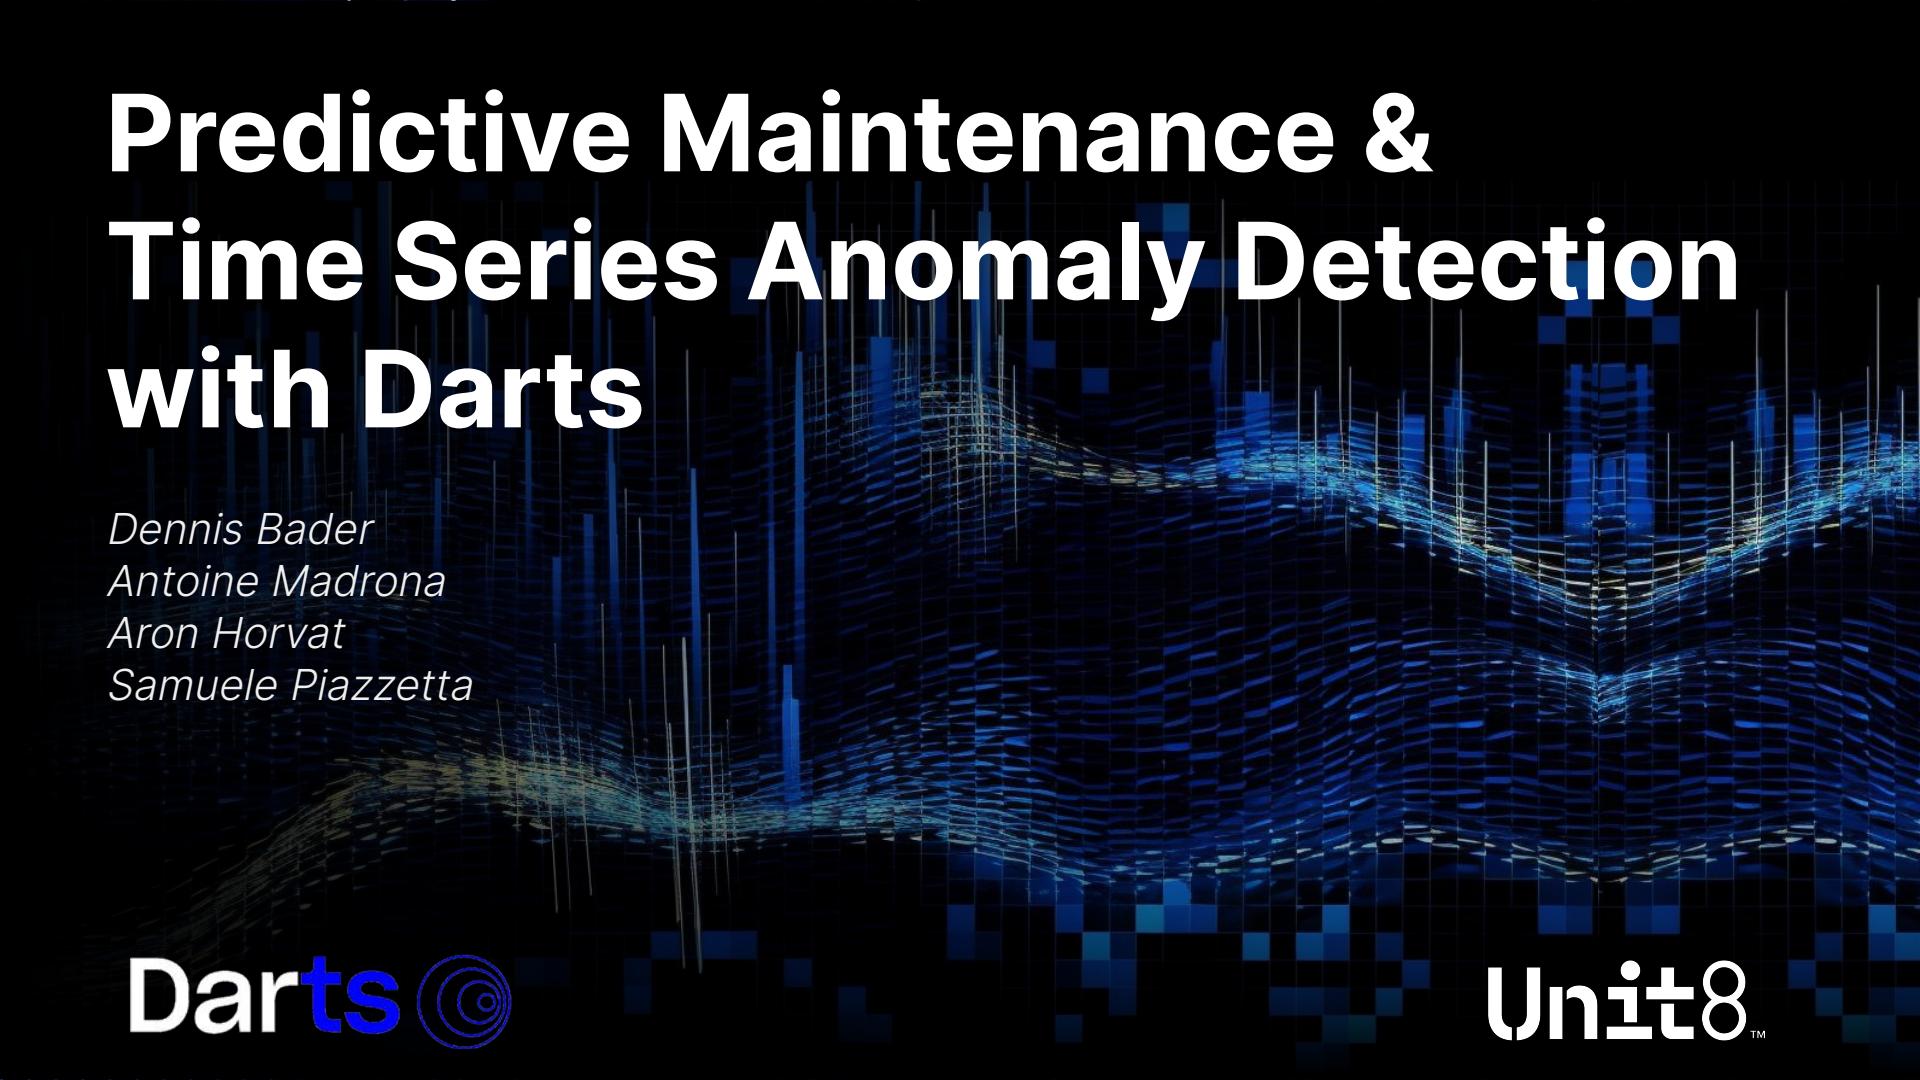


Predictive Maintenance & Time Series Anomaly Detection with Darts



*Dennis Bader
Antoine Madrona
Aron Horvat
Samuele Piazzetta*



Your Unit8 team today



Dennis Bader

Data Scientist &
Darts Lead Developer

MSc ETH
Mechanical Engineering



Antoine Madrona

Data Scientist &
Darts Developer

MSc EPFL
Life Science Engineering



Aron Horvath

Data Scientist &
Data Engineer

PostDoc University Zurich
Clinical Data Science



**Samuele Giuliano
Piazzetta**

Data Scientist &
Darts Developer

MSc ETH
Computer Science



Unit8 in a nutshell

Turning data into value

Unit8 is a leading Swiss data services company with a mission to help non-digital native companies transform data into value with a mix of data science, analytics, and AI.

Selected Clients



DAIMLER

MERCK

Firmenich

Swiss Re

helvetia

UV

TECAN.

Unit8

unit8.com

Key facts:

- Established in 2017
- 110+ employees and growing
- 5 locations: Zürich, Lausanne, Bern (Switzerland), Frankfurt (Germany) and Krakow, Wroclaw (Poland)
- 160+ data & AI projects completed

We operate at the intersection of technology and business to accompany our customers at every step of their data & AI journey

Services portfolio:

- Data & AI Consulting
- Data & AI Solutions
- Data & AI Platforms
- Data & AI Operations
- Data & AI Academy

Key Partners



Agenda Today

13:40 - 14:10 Intro to Forecasting with Darts

14:10 - 14:30 <Example Notebook>

14:30 - 14:40 Intro to Anomaly Detection

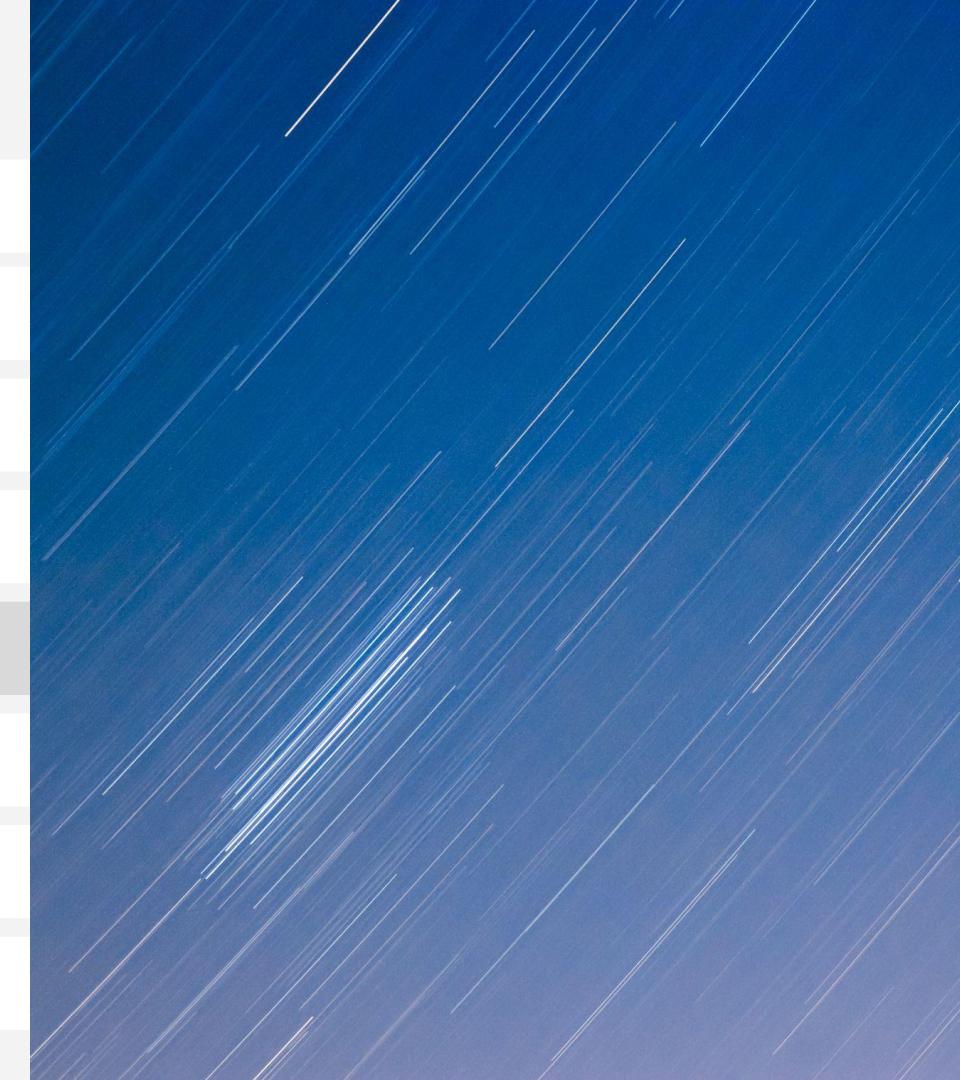
14:40 - 15:00 <Example Notebook>

15:00 - 15:30 Break

15:30 - 15:45 Intro to Predictive Maintenance

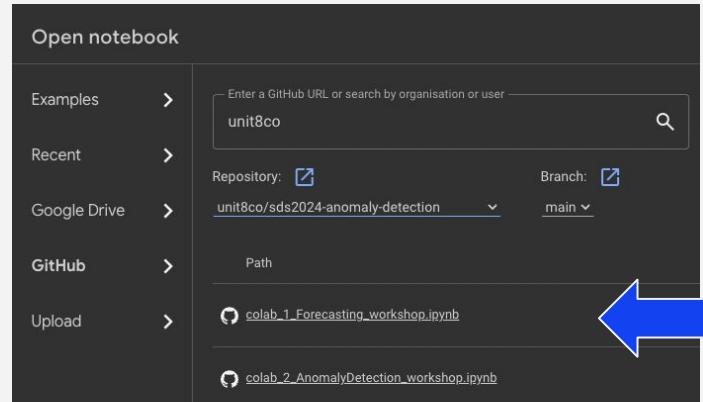
15:45 - 16:45 <Example Notebook>

16:45 - 17:00 Outro



About The Example Notebooks

- **3 example notebooks** to work on with **Google Colab**
 - Open: <https://colab.google/>
 - Click on “Open Colab”
 - On the left sidebar click GitHub
 - GitHub URL: “unit8co”
 - Repository: “unit8co/sds2024-anomaly-detection”
 - The three notebooks to work on are:
 - “colab_1_Forecasting_workshop.ipynb”
 - “colab_2_AnomalyDetection_workshop.ipynb”
 - “colab_3_PredictiveMaintenance_workshop.ipynb”
 - **Run up to 3 notebooks simultaneously** in different tabs
 - **Run the notebooks** until you receive the first error (first cells install dependencies and data)
- We also provide a **solution** for each example with prefix “colab_solution*”



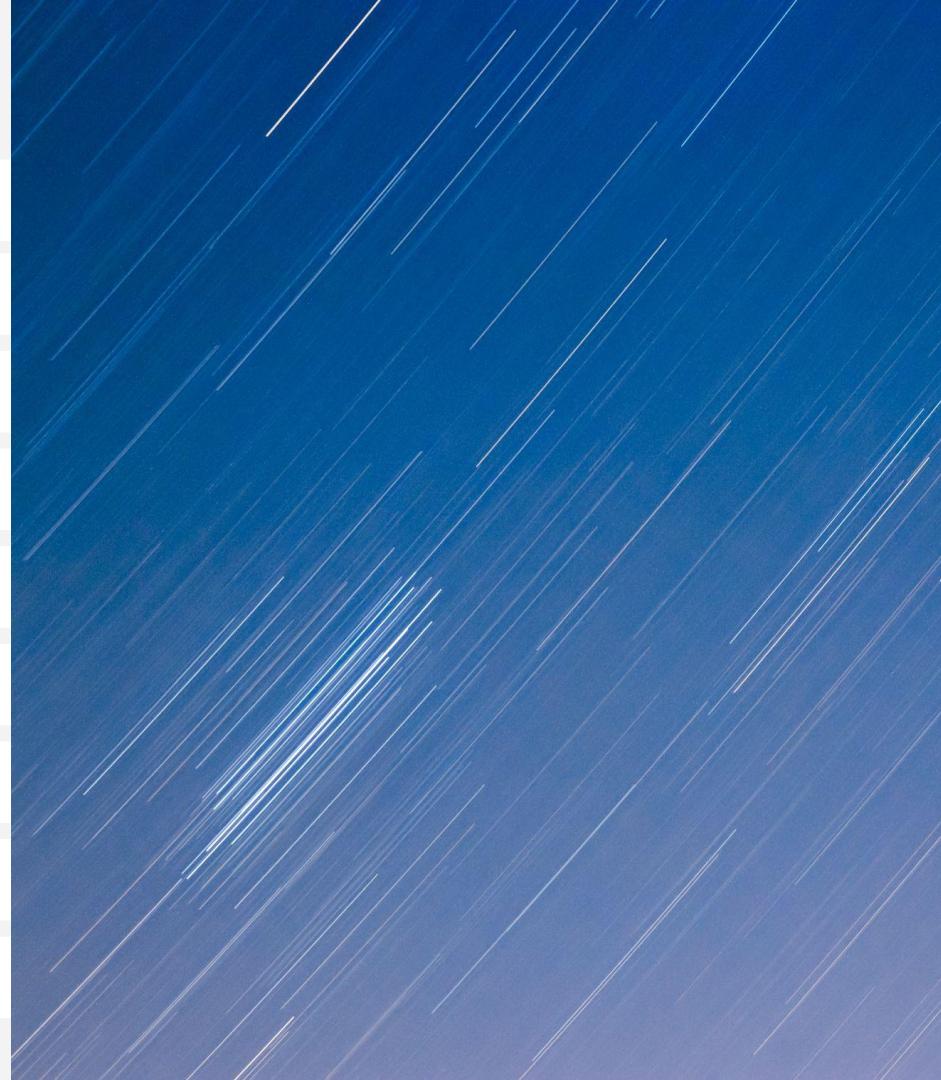
Timeseries Forecasting with Darts

Darts 

Unit⁸™

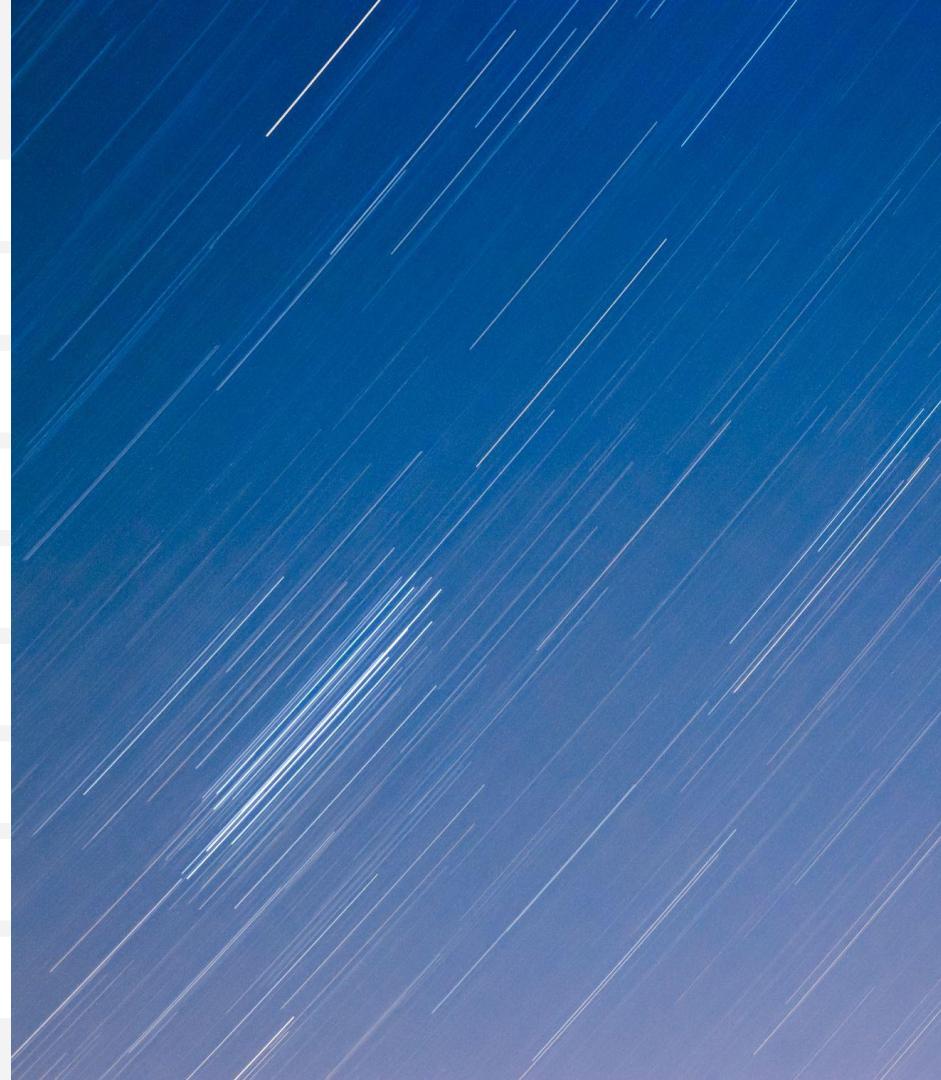
Agenda

- 1 Intro to Darts
- 2 TimeSeries Class
- 3 Forecasting models
- 4 Simple forecasting example
- 5 Historical forecasting and backtesting
- 6 Covariates
- 7 Probabilistic Forecasting
- 8 Training / predicting with multiple series
- 9 Summary and what's coming next



Agenda

- 1 Intro to Darts**
- 2 TimeSeries Class
- 3 Forecasting models
- 4 Simple forecasting example
- 5 Historical forecasting and backtesting
- 6 Covariates
- 7 Probabilistic Forecasting
- 8 Training / predicting with multiple series
- 9 Summary and what's coming next



Why Forecasting / Anomaly Detection?

Retail



How can we ensure that the consumer demand for every product is met without overspending and creating waste?

Energy



How much energy should we produce?

How much energy can we produce?

What will the electricity price be tomorrow?

Predictive Maintenance



How can we detect anomalous device behavior before serious damage occurs?



Darts

Darts is Unit8's Python library for
easy time series manipulation,
forecasting and anomaly detection.



Immutable `TimeSeries`
class as basic building
block.



Unified `fit()`, `predict()`
interface.



40+ classical models &
state-of-the-art ML
approaches for forecasting
and anomaly detection



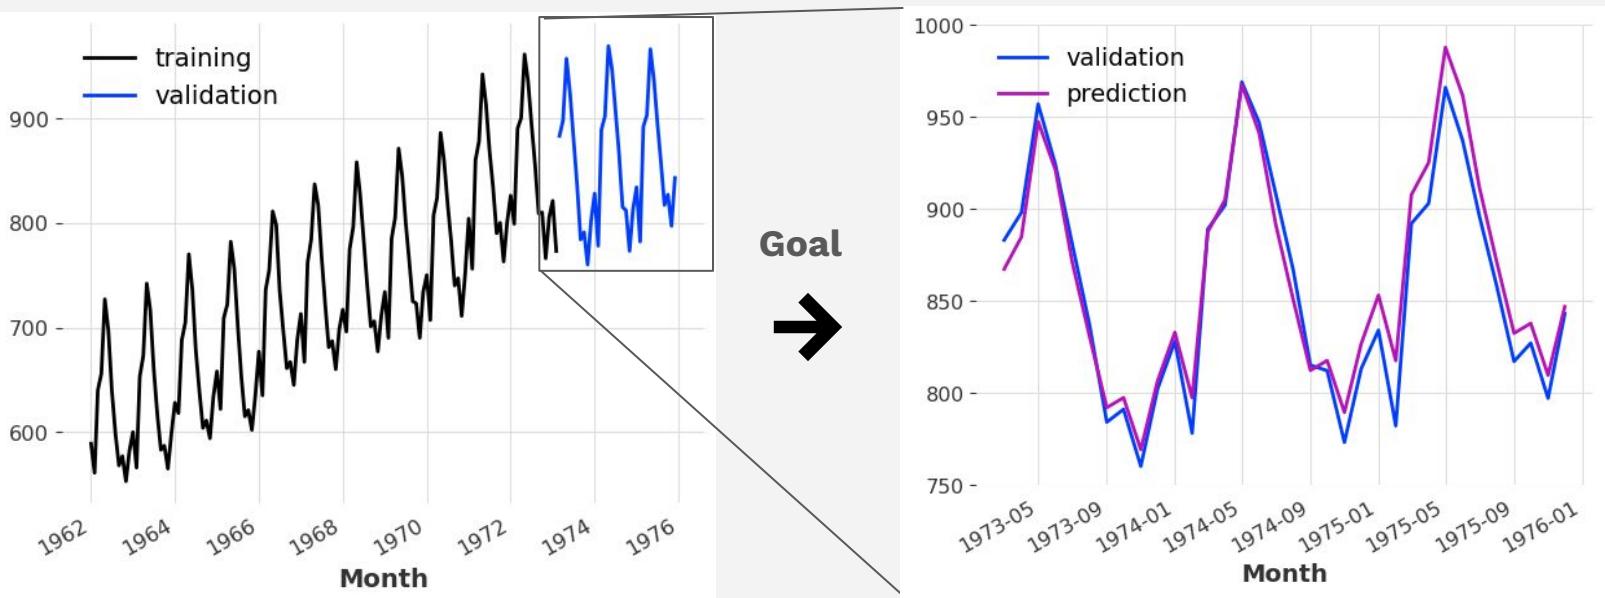
User-friendliness: intuitive
API and reasonable defaults.



Darts

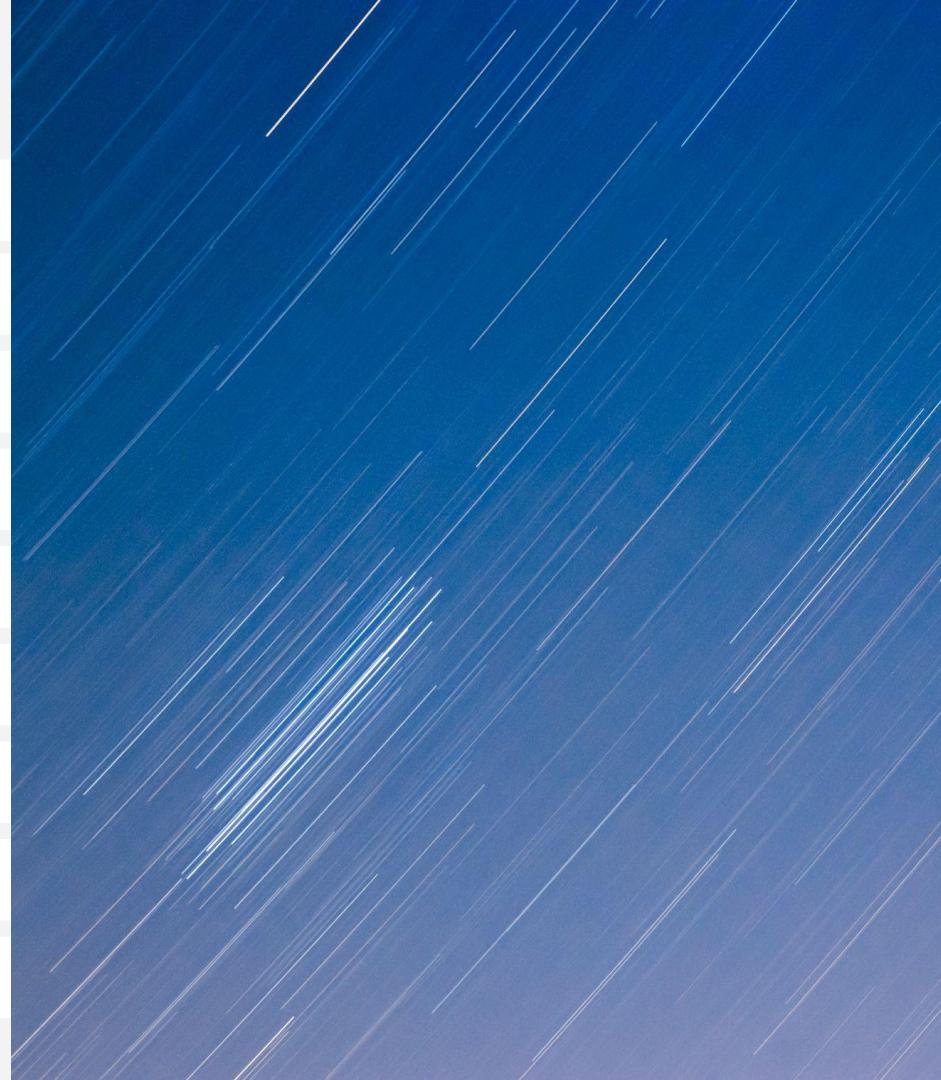


Try out 40+ models in only a few lines of code

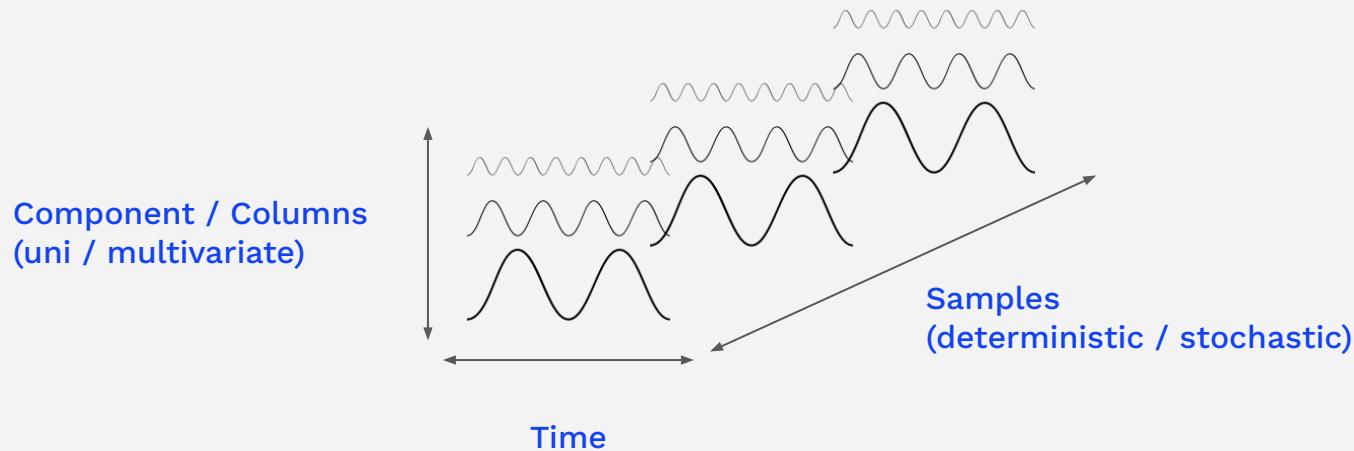


Agenda

- 1 Intro to Darts
- 2 TimeSeries Class**
- 3 Forecasting models
- 4 Simple forecasting example
- 5 Historical forecasting and backtesting
- 6 Covariates
- 7 Probabilistic Forecasting
- 8 Training / predicting with multiple series
- 9 Summary and what's coming next



TimeSeries: Darts' Container for storing and manipulating time series data



The TimeSeries object

```
from darts import TimeSeries  
  
# create time series  
series = TimeSeries.from_csv('milk.csv', time_col='Month')  
series.plot()
```

Factory Method Support

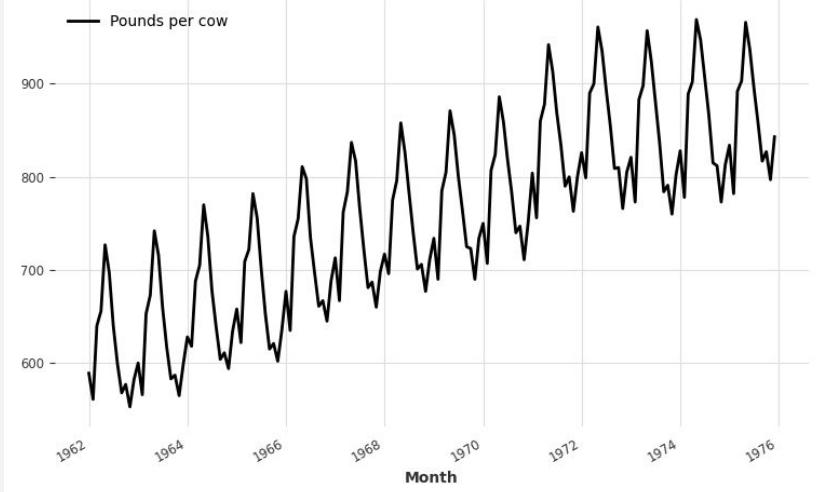


pandas



xarray

& other
file formats



TimeSeries Manipulation

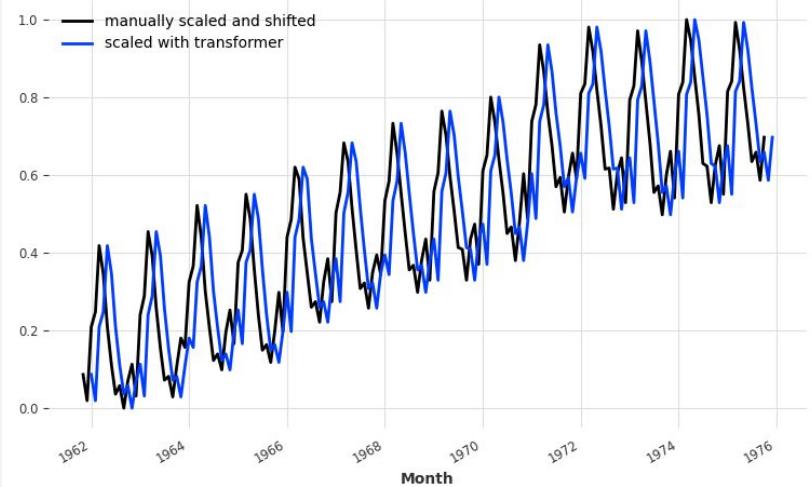
- **Methods and Arithmetic Operations**

```
# manually scale and shift  
((series - 553.0) / 416.0).shift(-2).plot()
```

- **Data processing module**

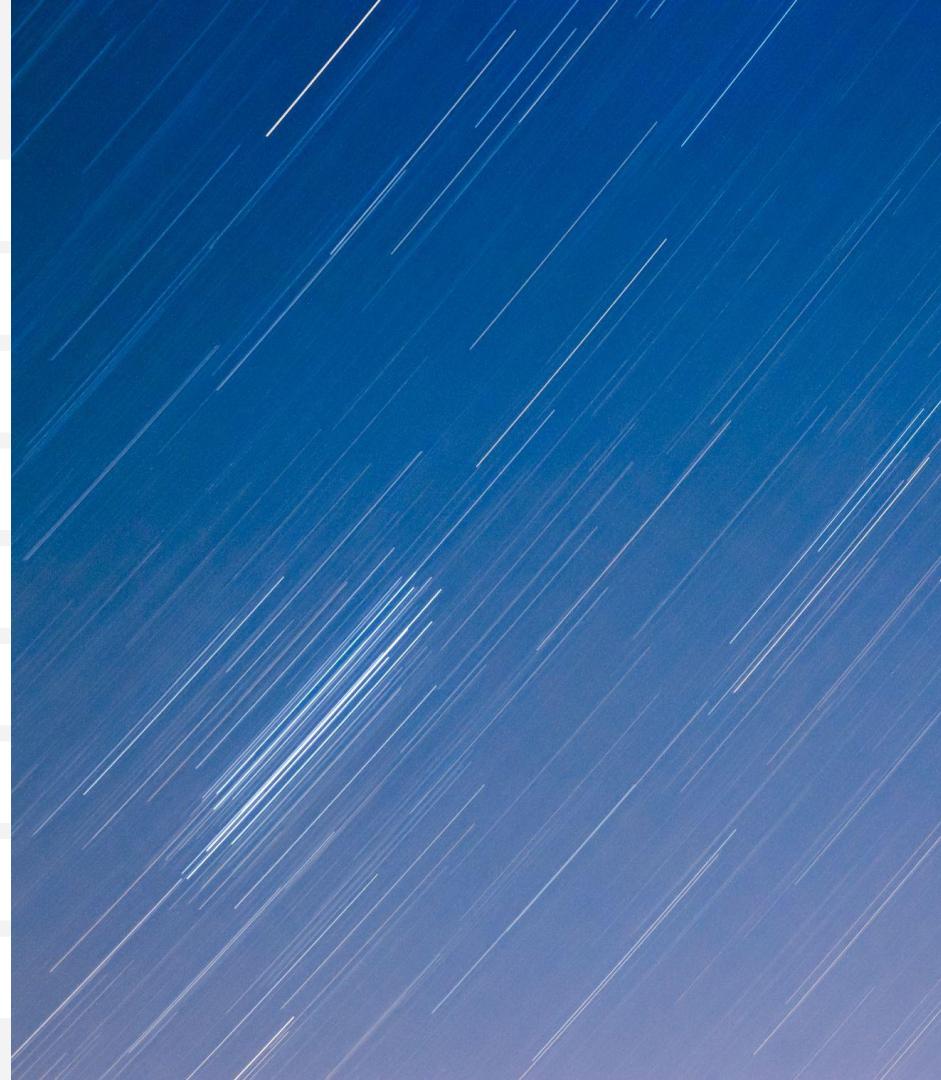
- Transformers
- Missing value handling
- Windowing

```
from darts.dataprocessing.transformers import Scaler  
  
# scale with a data transformer  
min_max_scaler = Scaler()  
min_max_scaler.fit_transform(series).plot()
```



Agenda

- 1 Intro to Darts
- 2 TimeSeries Class
- 3 Forecasting models**
- 4 Simple forecasting example
- 5 Historical forecasting and backtesting
- 6 Covariates
- 7 Probabilistic Forecasting
- 8 Training / predicting with multiple series
- 9 Summary and what's coming next



Darts ForecastingModels

Darts has a large collection of forecasting models; from statistical to deep learning models.

All models support:

- Univariate forecasting
- Unified sklearn-like `fit()`/`predict()` API

Additional support depending on the model:

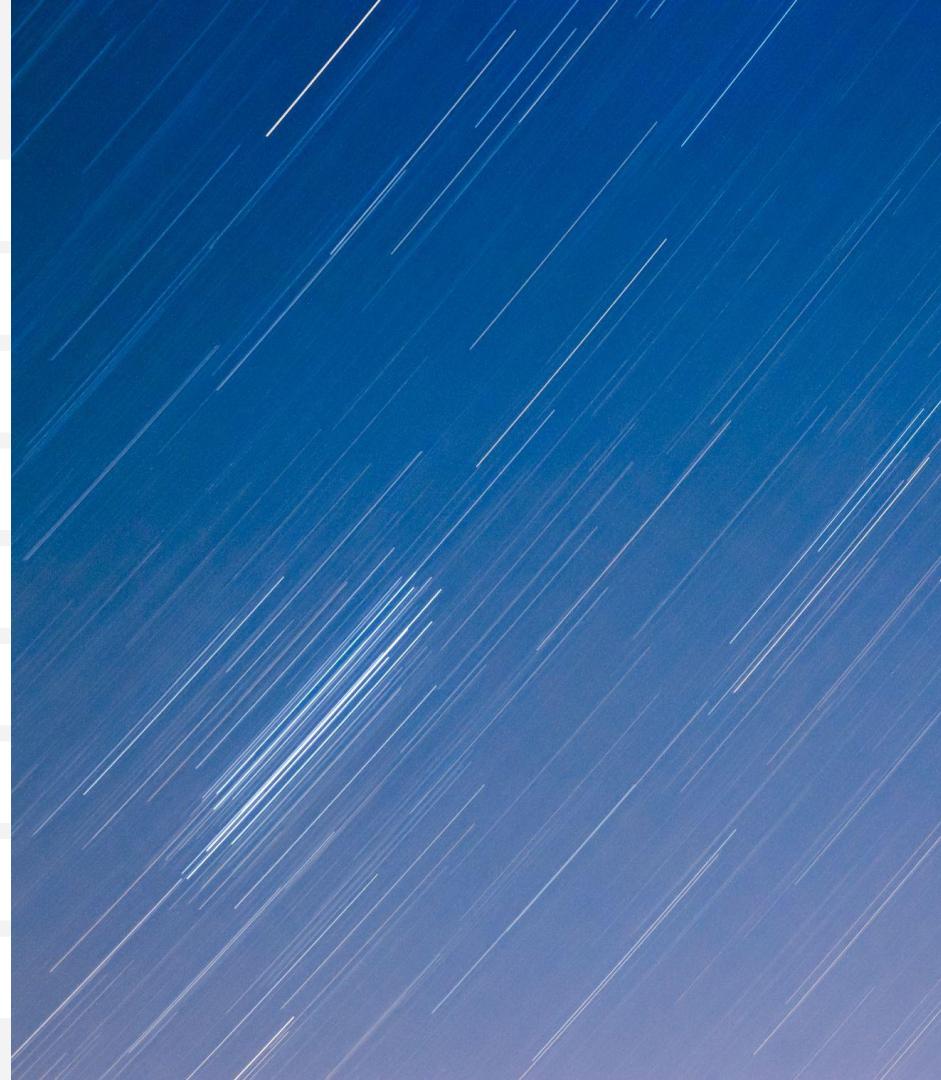
- Multivariate forecasting
- Training on multiple series
- Using covariates (external data) as input
- Probabilistic forecasting
- Neural-network related features

Model	Sources	Covariates Support:		Probabilistic Forecasting:	Training & Forecasting on Multiple Series
		Target Series Support:	Past-observed/Future-known/Static		
		Univariate/Multivariate	Sampled/Distribution Parameters		
Baseline Models (LocalForecastingModel)					
NaiveMean		■ ■	■ ■ ■ ■	■ ■	■
NaiveSeasonal		■ ■	■ ■ ■ ■	■ ■	■
NaiveDrift		■ ■	■ ■ ■ ■	■ ■	■
NaiveMovingAverage		■ ■	■ ■ ■ ■	■ ■	■
Statistical / Classic Models (LocalForecastingModel)					
ARIMA		■ ■	■ ■ ■ ■	■ ■	■
VARIMA		■ ■	■ ■ ■ ■	■ ■	■
AutoARIMA		■ ■	■ ■ ■ ■	■ ■	■
StatsForecastAutoArima (faster AutoARIMA)	Nixtla's statsforecast	■ ■	■ ■ ■ ■	■ ■	■

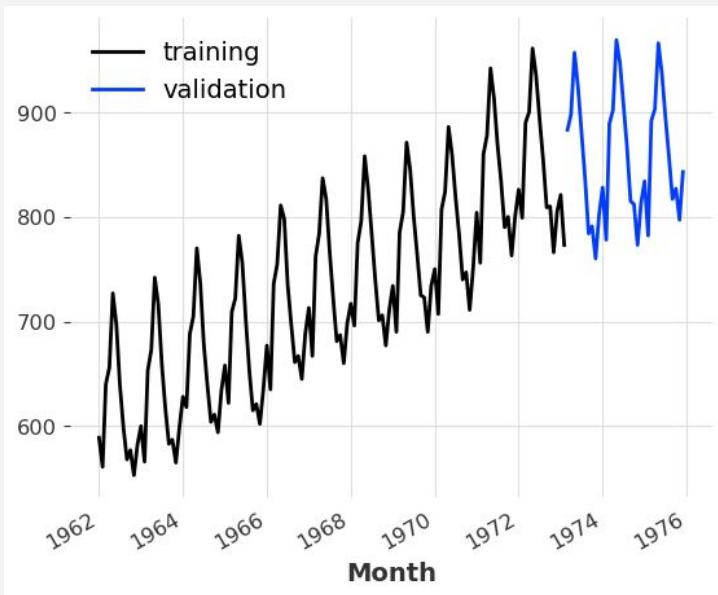
First 8 out of 40+ Forecasting Models, see
<https://unit8co.github.io/darts/>

Agenda

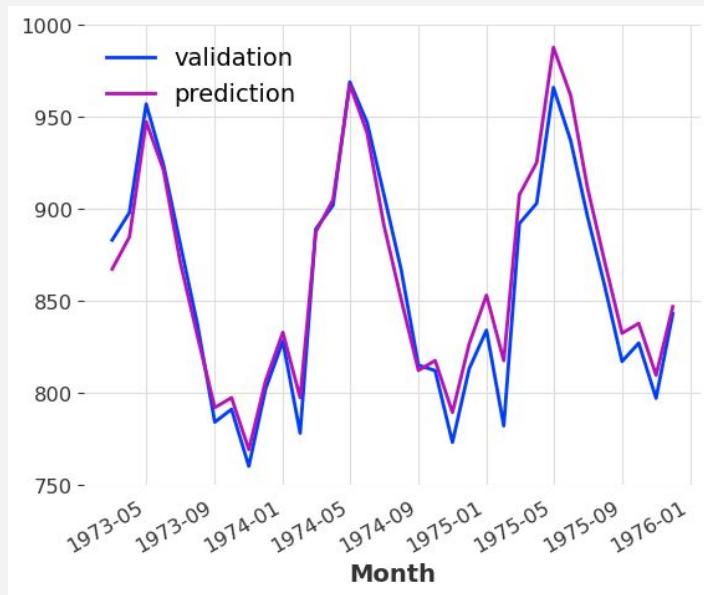
- 1 Intro to Darts
- 2 TimeSeries Class
- 3 Forecasting models
- 4 Simple forecasting example**
- 5 Historical forecasting and backtesting
- 6 Covariates
- 7 Probabilistic Forecasting
- 8 Training / predicting with multiple series
- 9 Summary and what's coming next



Goal



*Accurate
Forecasts*



TimeSeries

Processing

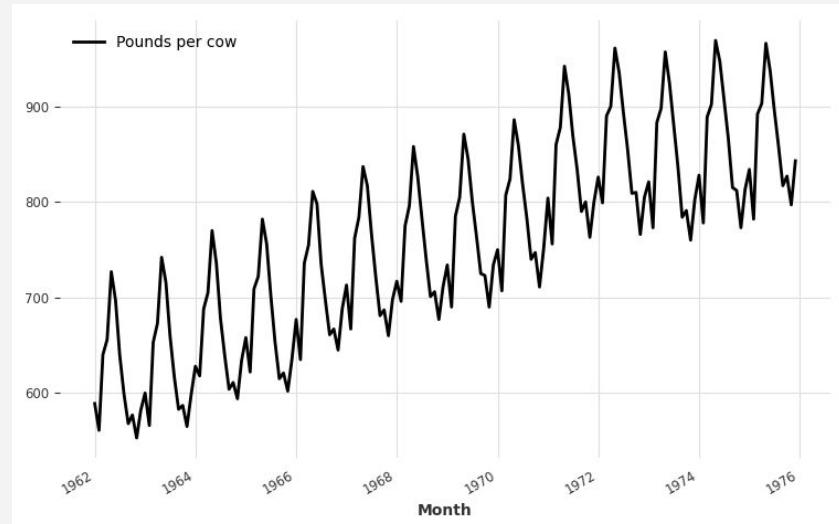
Forecasting

Evaluating

The TimeSeries object

```
from darts import TimeSeries

# create time series
series = TimeSeries.from_csv('milk.csv', time_col='Month')
series.plot()
```



TimeSeries

Processing

Forecasting

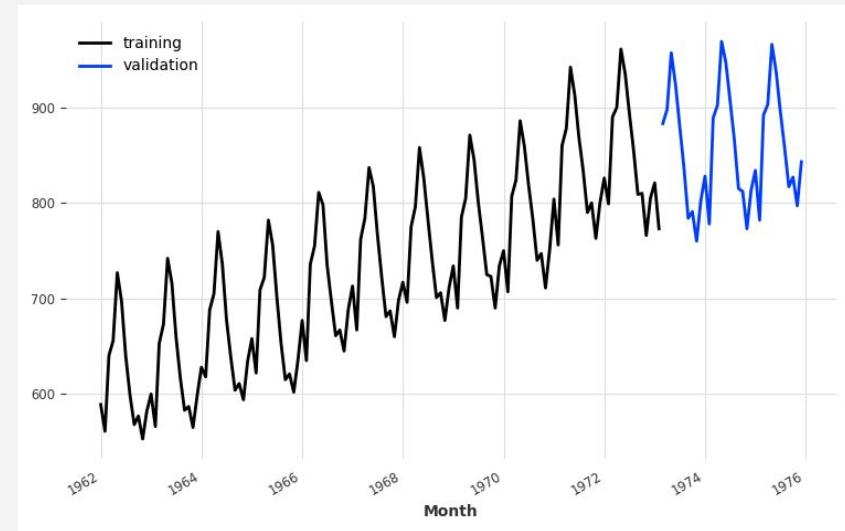
Evaluating

Training / validation split

```
from darts import TimeSeries

# create time series
series = TimeSeries.from_csv('milk.csv', time_col='Month')

# split train/val
train, test = series.split_after(0.8)
```



TimeSeries

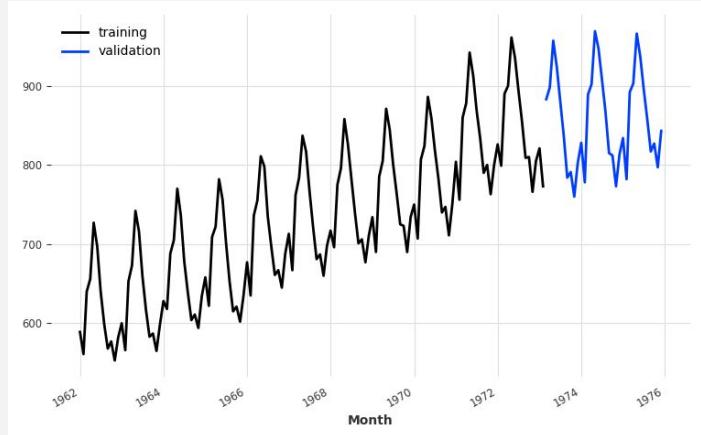
Processing

Forecasting

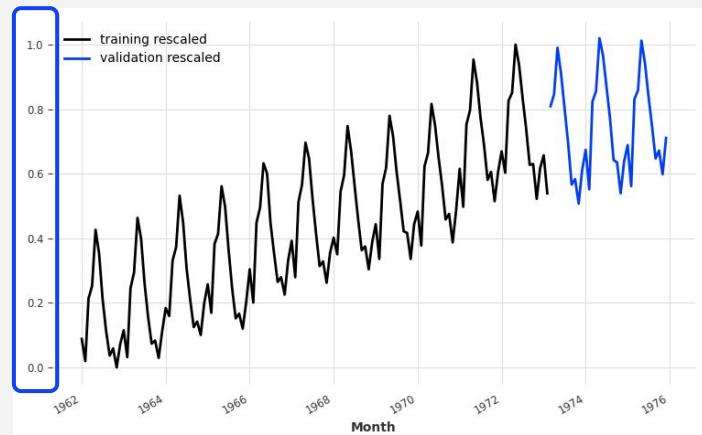
Evaluating

Normalizing TimeSeries

```
from darts.dataprocessing.transformers import Scaler  
  
scaler = Scaler()  
scaler.fit(series) # learn min and max  
rescaled = scaler.transform(series)
```

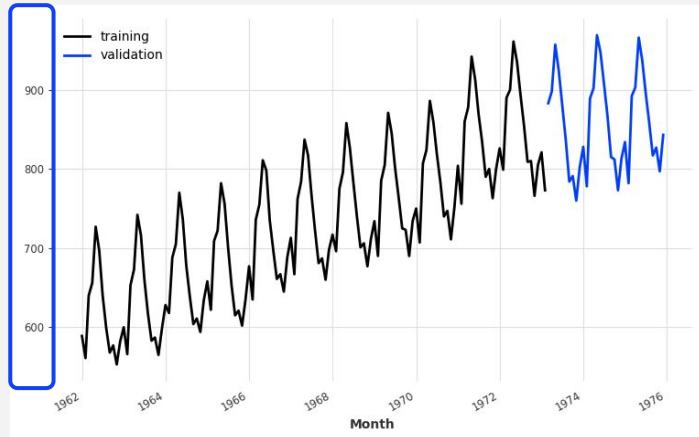


Normalizing

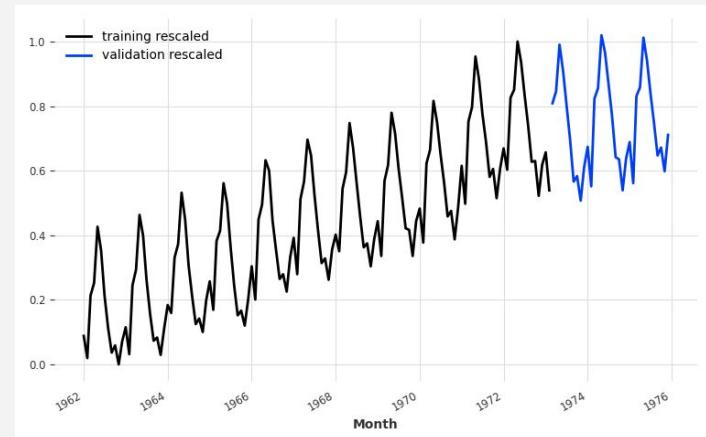


Normalizing TimeSeries

```
original_series = scaler.inverse_transform(rescaled)
```

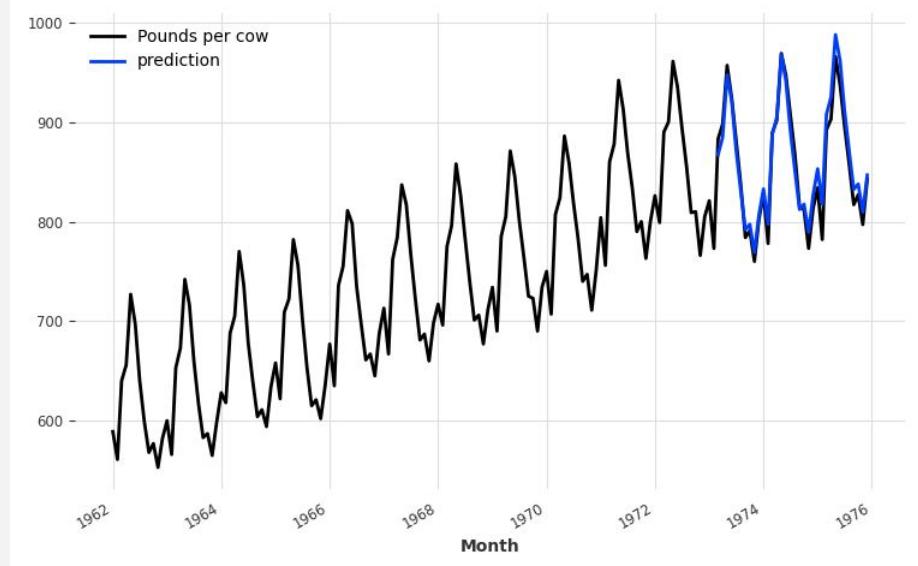


*Re-scaling
back*



Forecasting – Exponential Smoothing

```
from darts.models import ExponentialSmoothing  
  
model = ExponentialSmoothing()  
model.fit(train)  
pred = model.predict(n=len(test))
```



TimeSeries

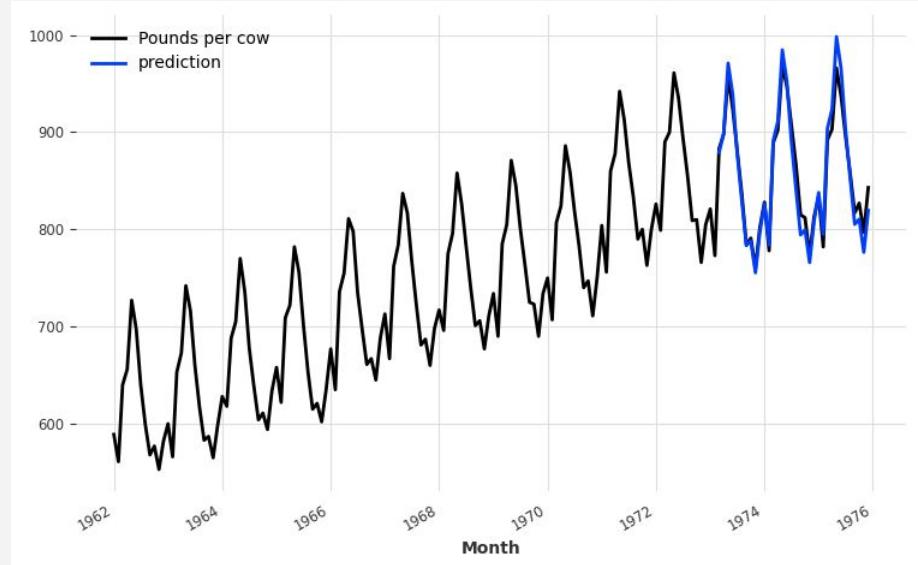
Processing

Forecasting

Evaluating

Forecasting – Theta

```
from darts.models import Theta  
  
model = Theta()  
model.fit(train)  
pred = model.predict(n=len(test))
```



TimeSeries

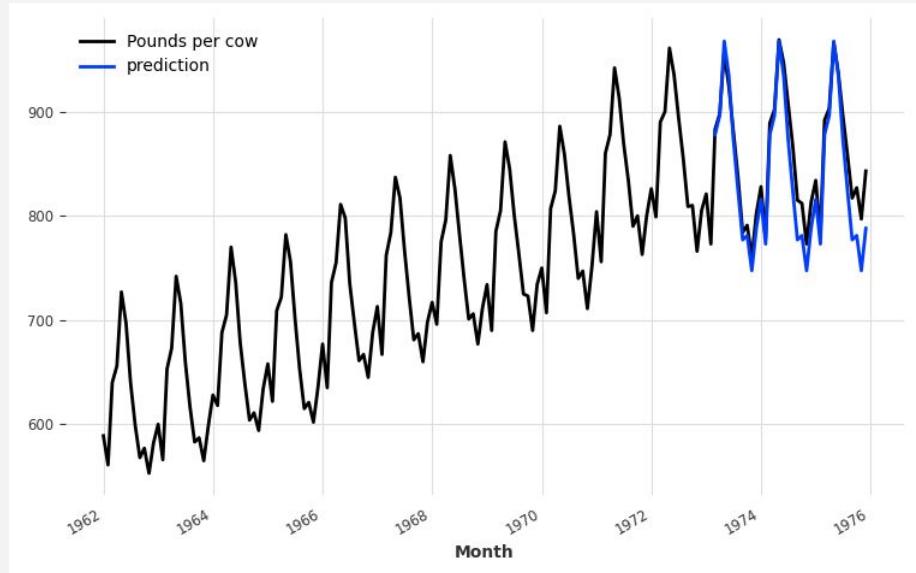
Processing

Forecasting

Evaluating

Specifying model parameters

```
from darts.models import Theta  
  
model = Theta(theta=1)  
model.fit(train)  
pred = model.predict(n=len(test))
```



TimeSeries

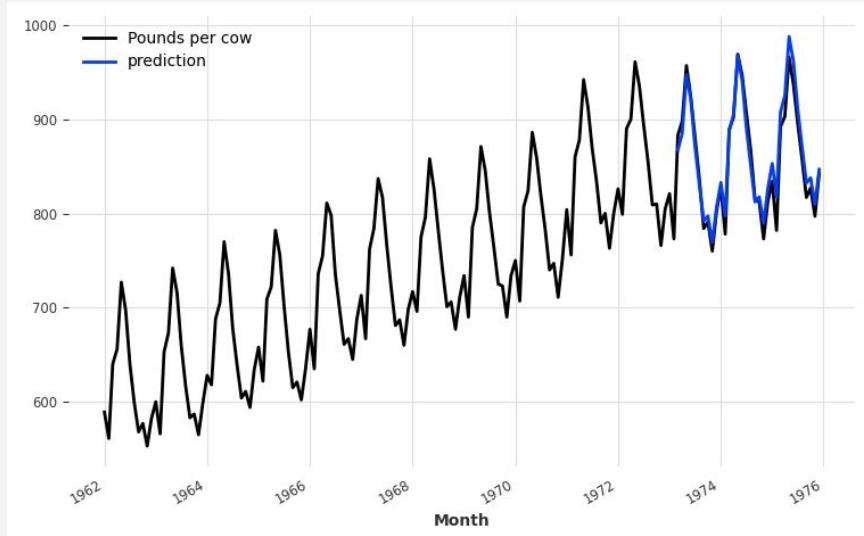
Processing

Forecasting

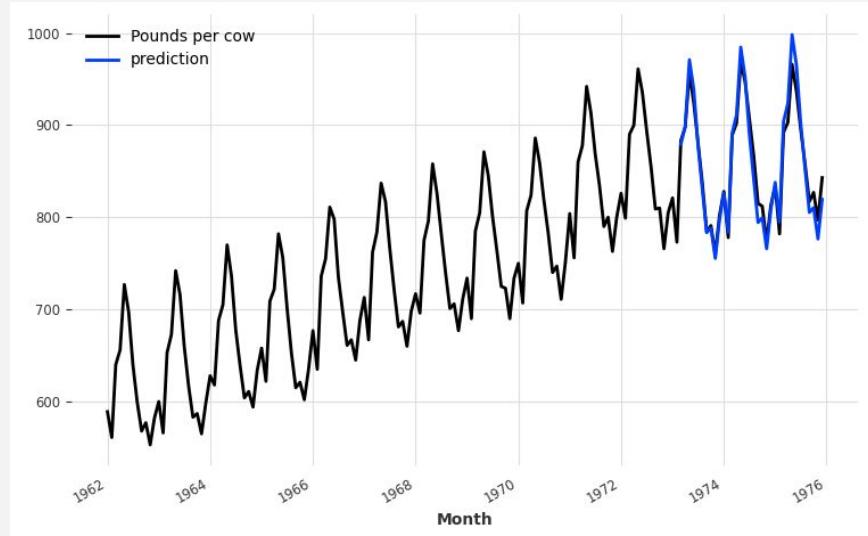
Evaluating

Evaluating predictions – Which one is better?

Exponential Smoothing



Theta



TimeSeries

Processing

Forecasting

Evaluating

Metrics

Many different scores can be computed – Darts lets you import the ones you need.

```
from darts.metrics import mape  
  
score_mape = mape(test, pred)
```

```
from darts.metrics import smape  
  
score_smape = smape(test, pred)
```

TimeSeries

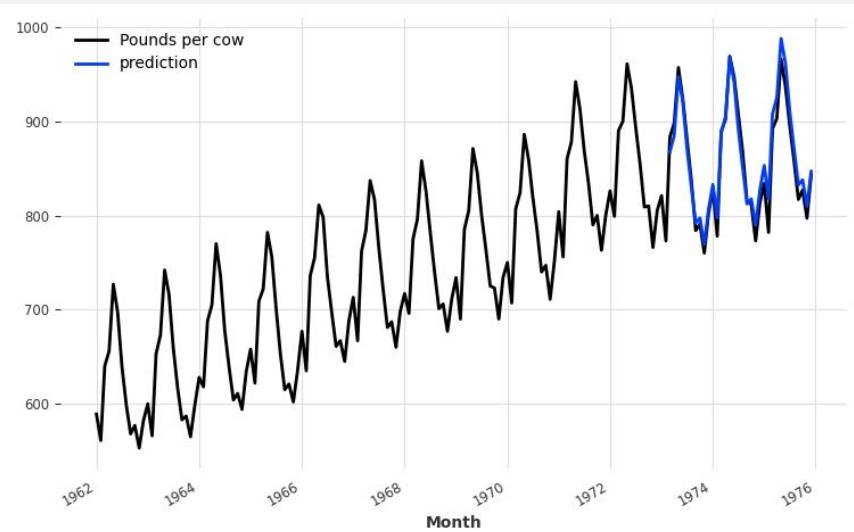
Processing

Forecasting

Evaluating

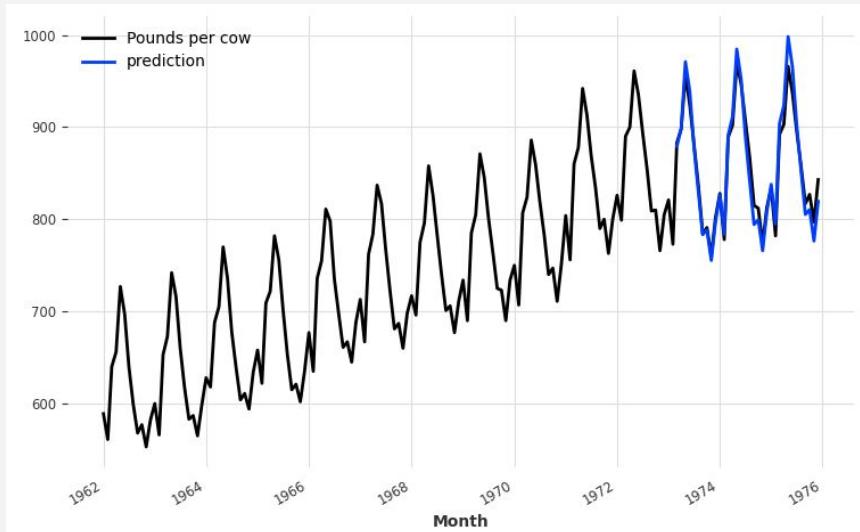
Which one is better?

Exponential Smoothing



sMAPE: ~1.39%

Theta



sMAPE: ~1.28% 

TimeSeries

Processing

Forecasting

Evaluating

Recap

End-to-end example

```
from darts import TimeSeries
from darts.dataprocessing.transformers import Scaler
from darts.metrics import smape
from darts.models import Theta

horizon = 12

# read searies, split train/val/test/...
series = TimeSeries.from_xxx(...)
train, test = series[:-horizon], series[-horizon:]

# scale train/val/...
scaler = Scaler()
train = scaler.fit_transform(train)
val = scaler.transform(test)

# build model with some parameters
model = Theta(...)

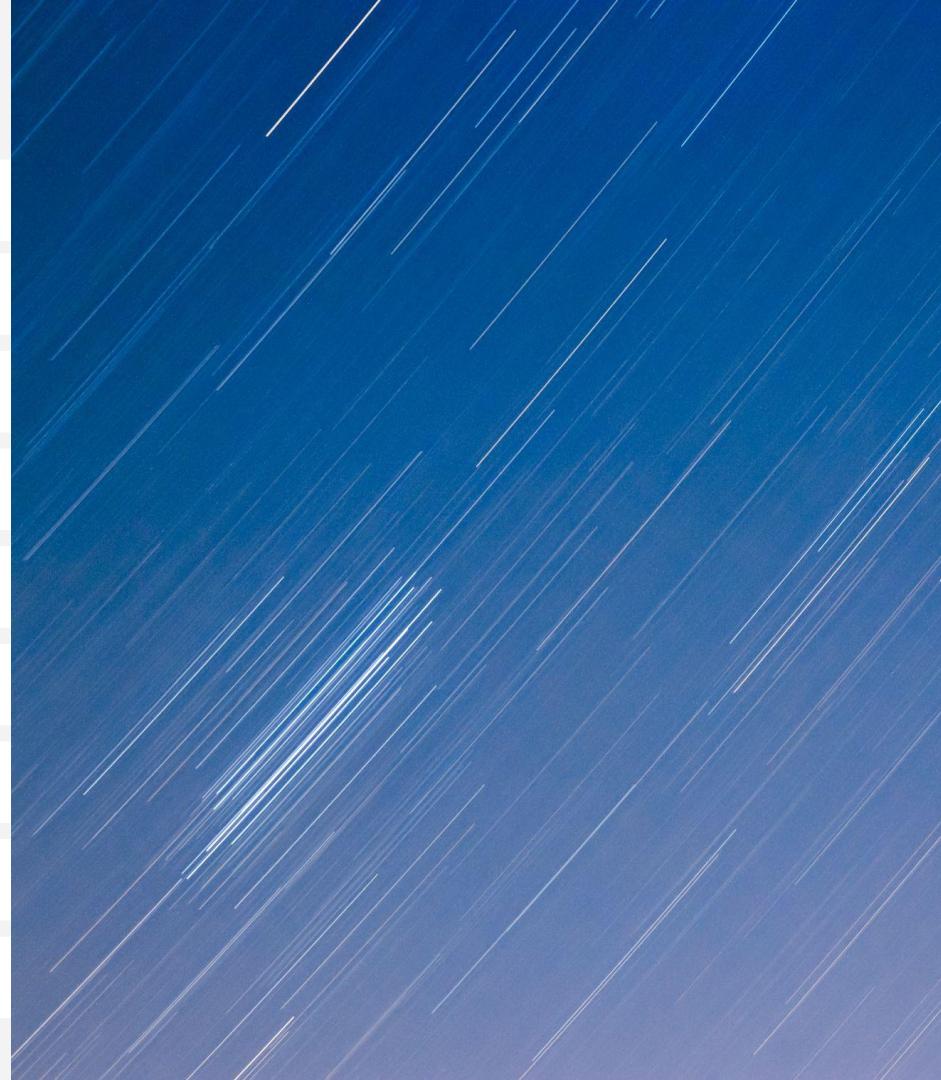
# train model
model.fit(train)

# get a prediction
pred: TimeSeries = model.predict(n=horizon)

# compute metrics
smape_test = smape(test, pred)
```

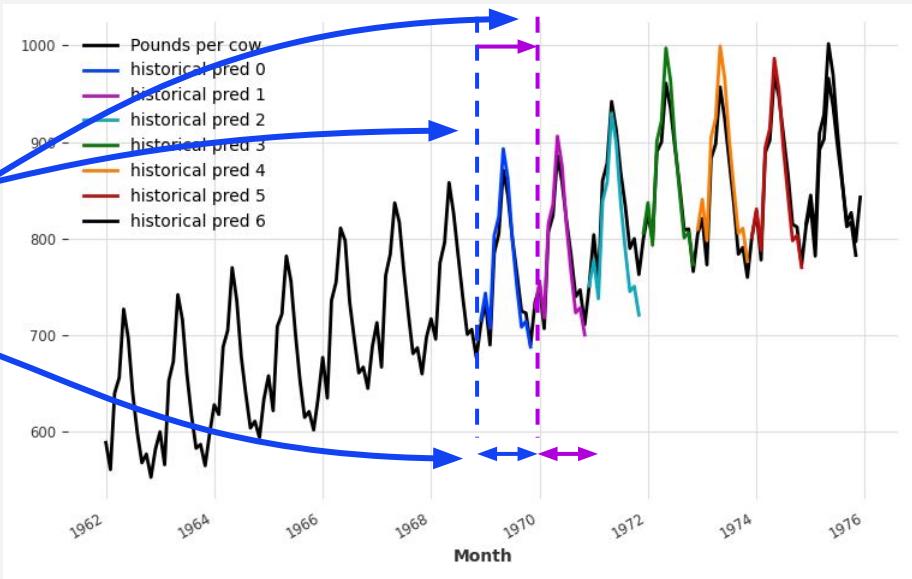
Agenda

- 1 Intro to Darts
- 2 TimeSeries Class
- 3 Forecasting models
- 4 Simple forecasting example
- 5 Historical forecasting and backtesting**
- 6 Training / predicting with multiple series
- 7 Covariates
- 8 Probabilistic Forecasting
- 9 Summary and what's coming next



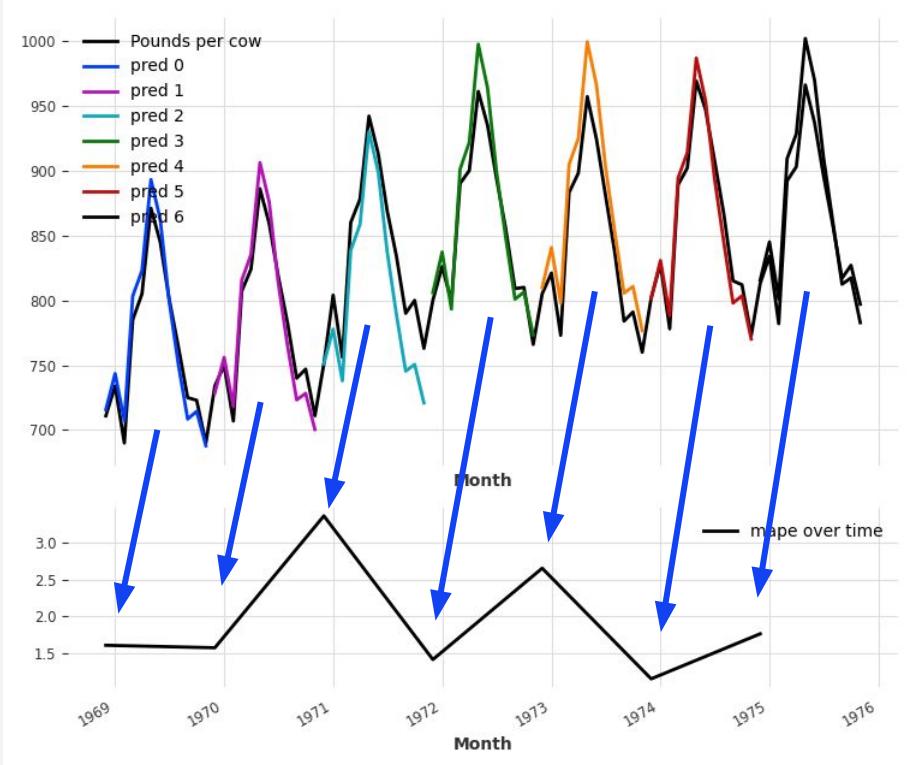
Historical simulations using Darts

```
forecasts = model.historical_forecasts(  
    series=series,  
    start=0.5,  
    forecast_horizon=12,  
    stride=12,  
    last_points_only=False,  
)
```



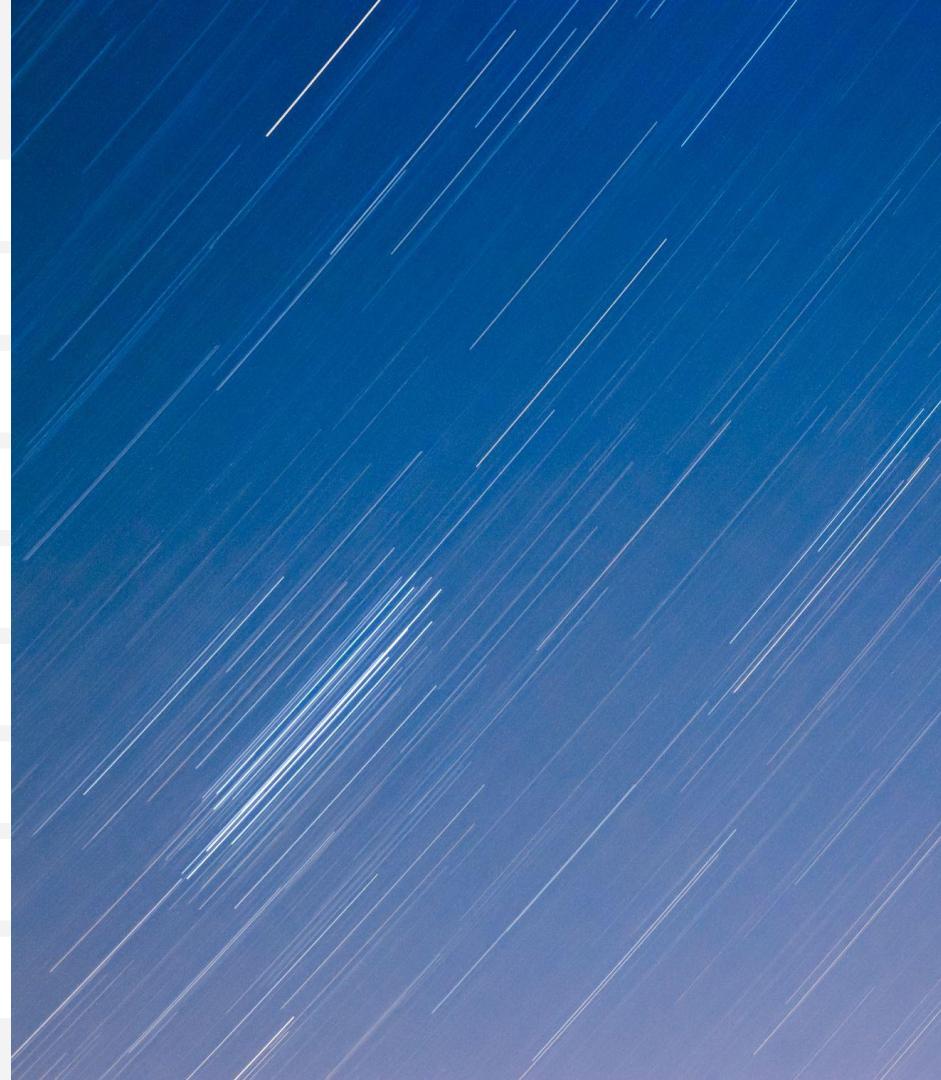
Backtesting models on historical data

```
backtest_errors = model.backtest(  
    historical_forecasts=forecasts,  
    series=series,  
    metric=mape,  
    reduction=None,    # np.mean, ...  
)
```



Agenda

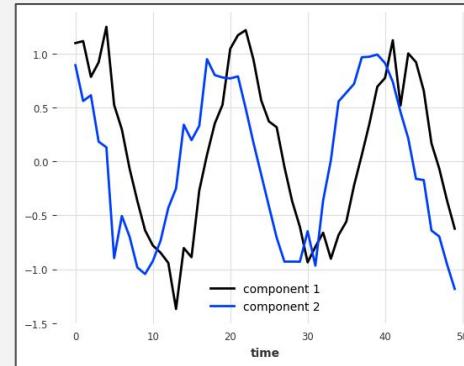
- 1 Intro to Darts
- 2 TimeSeries Class
- 3 Forecasting models
- 4 Simple forecasting example
- 5 Historical forecasting and backtesting
- 6 Training / predicting with multiple series**
- 7 Covariates
- 8 Probabilistic Forecasting
- 9 Summary and what's coming next



Multiple series means several uni- or multivariate series

In Darts, one `TimeSeries` can be:

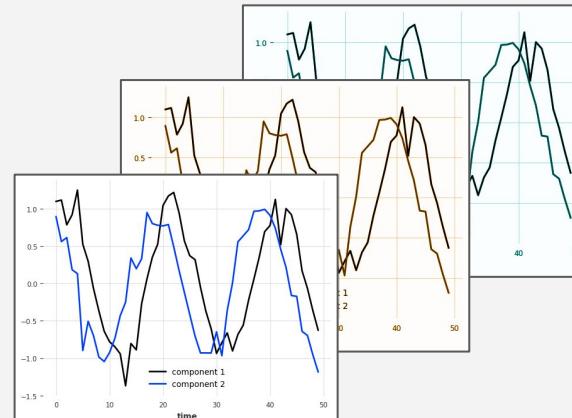
- **Univariate**: one column
- **Multivariate**: multiple columns



← one multivariate series
with two columns

Multiple series are several `TimeSeries` that:

- must have the same number of columns
- can have different time axes, scales,
trends, ...



← multiple multivariate
series

Back to the forecasting techniques

Local Forecasting Models

ARIMA/ VARIMA	Exponential Smoothing	Fast Fourier Transform
Theta	BATS/TBATS	Croston
Prophet	Baseline Models (Moving Average, Drift, Seasonal, ...)	...

Global Forecasting Models

Baseline Models (Fixed & Moving Aggregate/Mean, Drift, Seasonal)
Regression Models (scikit-learn, XGBoost, LightGBM, Catboost)
Neural Networks (RNN, NBEATS/NHiTS, TCN, Transformer, TFT, NLinear/DLinear, TiDE)
Ensemble Models* (Naive, Regression-based)

- Can only be trained on a **single** target series
- Tend to be simpler statistical or naive models
- Can predict $n \rightarrow \infty$

- Can be trained on **single** or **multiple** target series
- ML based-, and ensemble models
- Train on fixed-length sub-samples of input data

* depending on which forecasting models are ensembled

In Practice

Fitting and forecasting on collections of series

Single Series

```
from darts.models import SomeForecastingModel as Model

# create a series
series: TimeSeries = series1

model = Model(...)
model.fit(series)

# generate forecasts for a single series
preds: TimeSeries = model.predict(
    n=horizon,
)
```

Multiple Series

```
from darts.models import SomeGlobalModel as Model

# create a sequence of series
series: Sequence[TimeSeries] = [series1, series2]

model = Model(...)
model.fit(series)

# generate forecasts for each `series`
preds: List[TimeSeries] = model.predict(
    n=horizon,
    series=series, # or a completely new series
)
```

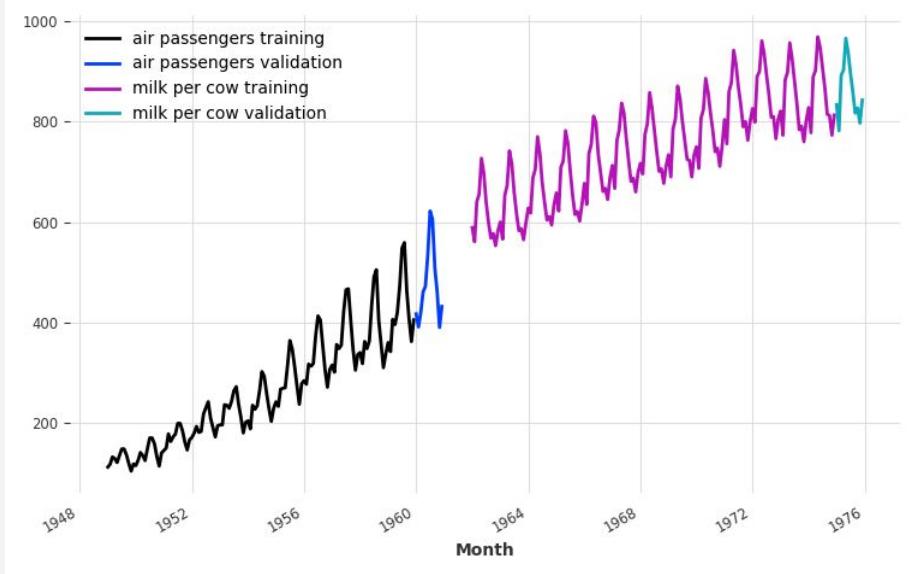
Looking at an example with [Multiple Series](#)

```
from darts.datasets import (
    MonthlyMilkDataset,
    AirPassengersDataset
)

# load some of Darts datasets for experimentation
milk = MonthlyMilkDataset().load()
psng = AirPassengersDataset().load()

# create train and test sets with multiple series
horizon = 12
psng_train, psng_test = psng[:-horizon], psng[-horizon:]
milk_train, milk_test = milk[:-horizon], milk[-horizon:]

train = [psng_train, milk_train]
test = [psng_test, milk_test]
```

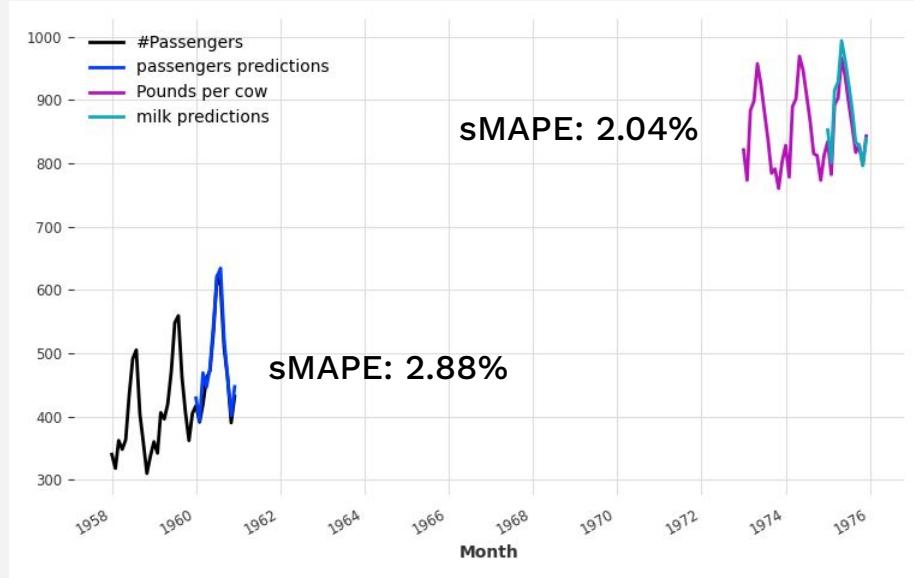


Training and forecasting on **Multiple Series** with a global model

```
# this is a neural network model
model = TiDEMModel(**tide_params)
model.fit(train)

# predict after the end of each series in `train`
preds = model.predict(n=horizon, series=train)

# compute metrics for each forecast
score = smape(test, preds)
```



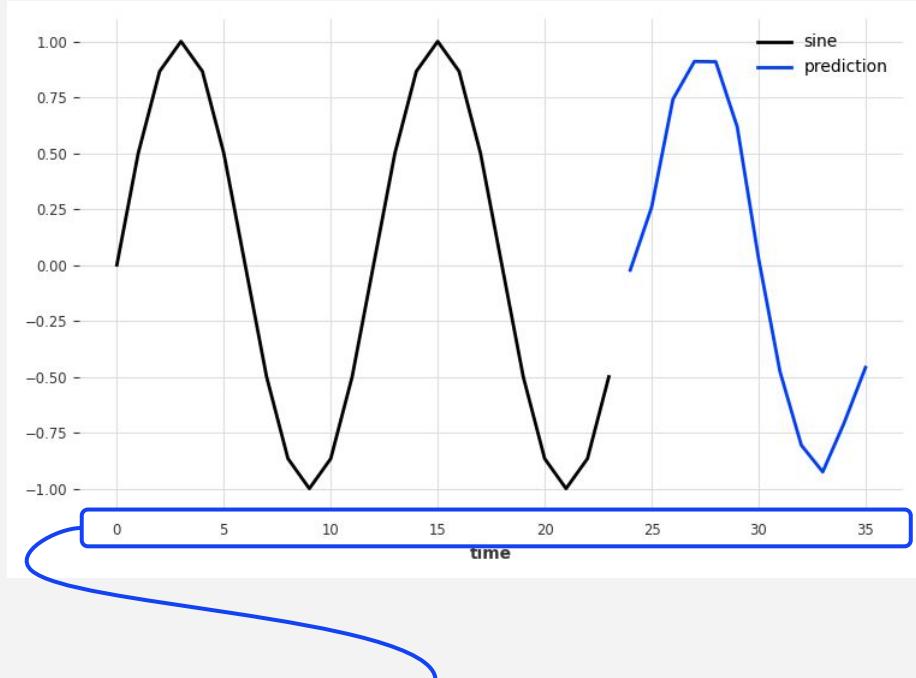
```
tide_params = {
    "input_chunk_length": 24,  # length of historic input per training sample
    "output_chunk_length": 12,  # length of future output per training sample
    "random_state": 0,  # for reproducibility
    "use_reversible_instance_norm": True,
}
```



Using the trained model to **forecast** an unobserved series

```
from darts.utils.timeseries_generation import sine_timeseries

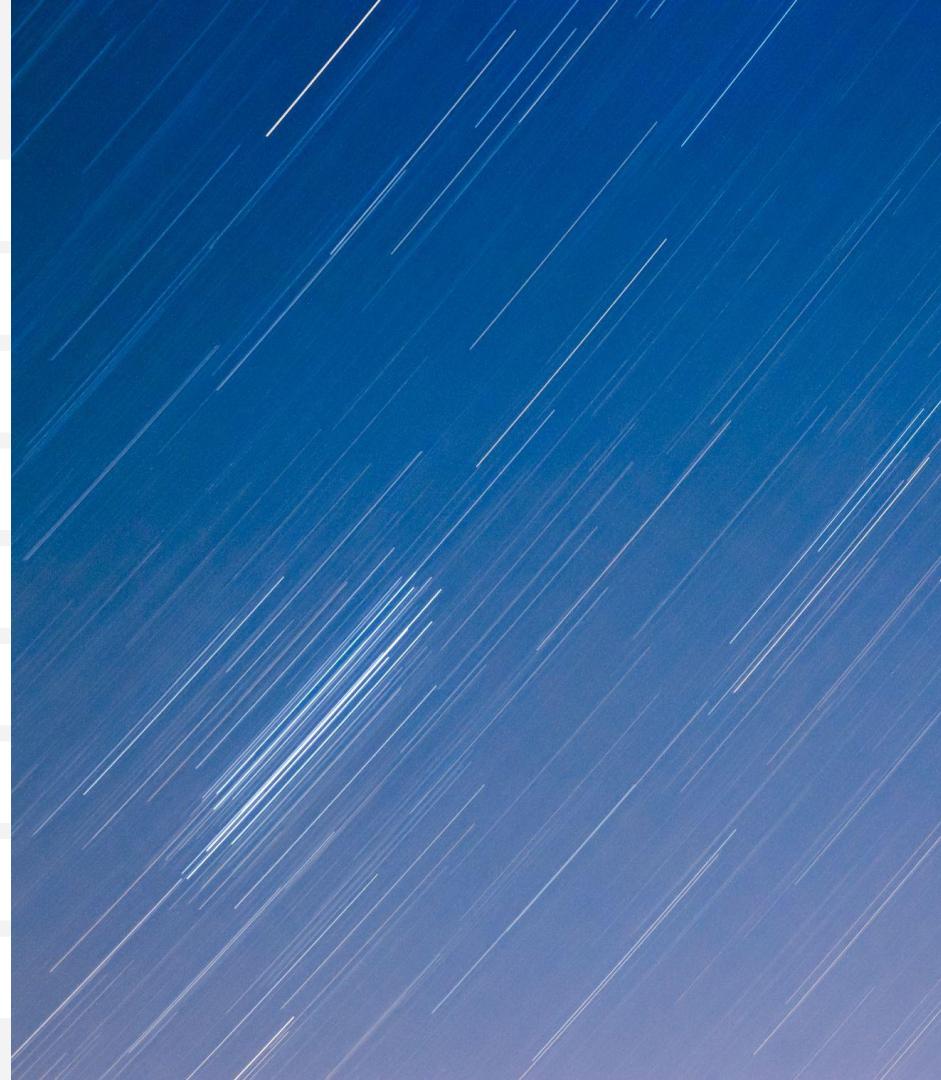
# generate a new sine curve series with similar seasonality
new_series = sine_timeseries(
    value_frequency=1/horizon,
    length=24,
    start=0 # give it an integer index
)
preds = model.predict(n=horizon, series=new_series)
```



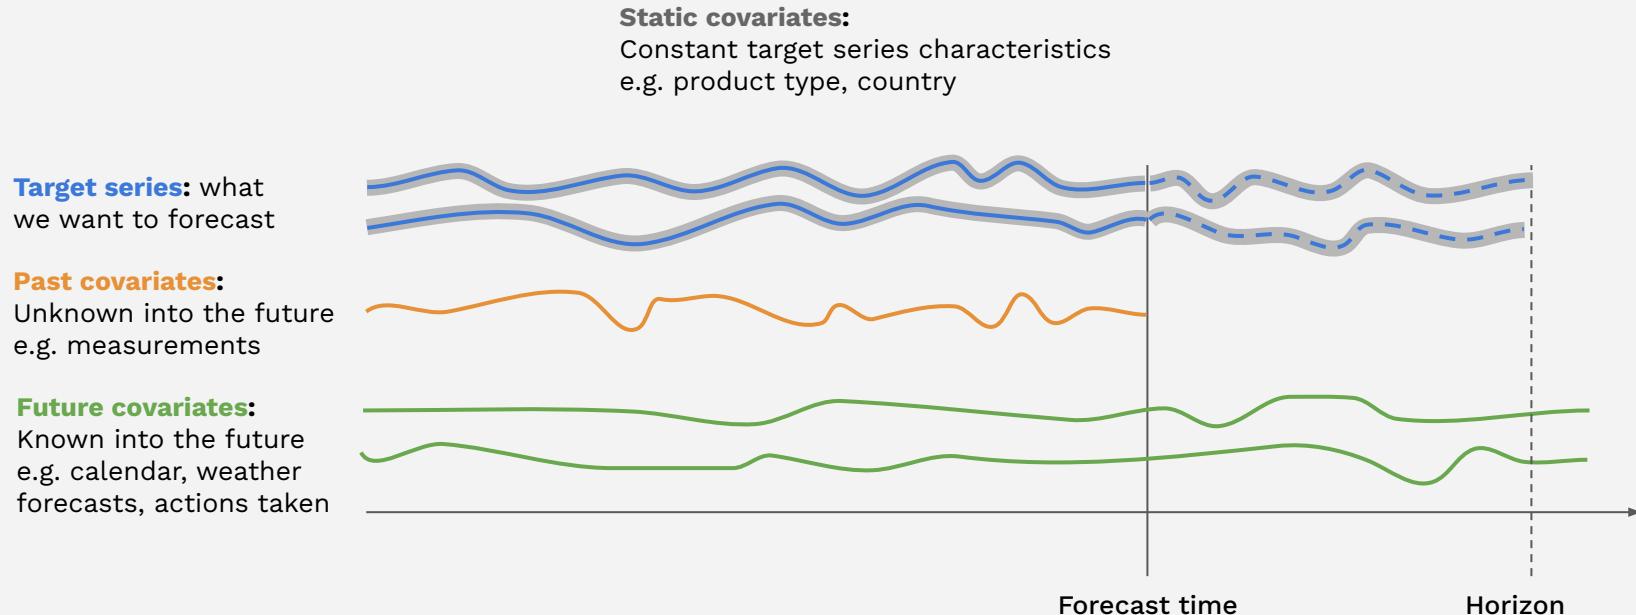
The time index and frequency doesn't have to match the one it was trained on (before: monthly time index, now: integer index)

Agenda

- 1 Intro to Darts
- 2 TimeSeries Class
- 3 Forecasting models
- 4 Simple forecasting example
- 5 Historical forecasting and backtesting
- 6 Training / predicting with multiple series
- 7 Covariates**
- 8 Probabilistic Forecasting
- 9 Summary and what's coming next



Including Past, Future & Static External Data



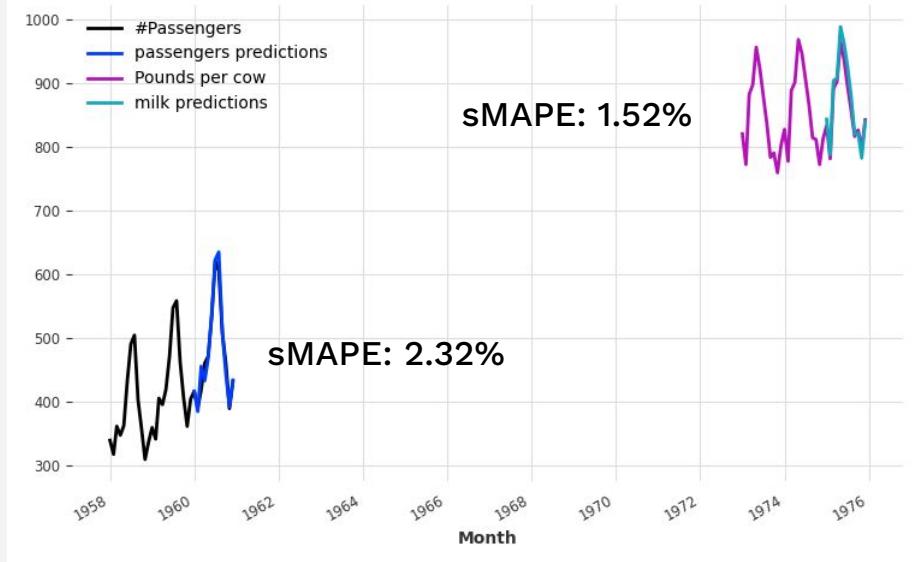
- `fit()` and `predict()` can accept `past_covariates` and/or `future_covariates`, depending on model
- Static covariates are embedded in the target series and can be used by some models
- Covariates can be uni- or multivariate
- Using `multiple target` series requires the same number of `past / future covariate series`
- Alignment of covariates with target is automatic



Now we train and forecast using the covariates

```
# let the model use static information
model = TiDEModel(
    **tide_params,
    use_static_covariates=True
)
# pass the covariates to fit() & predict()
model.fit(
    train,
    future_covariates=covariates,
)

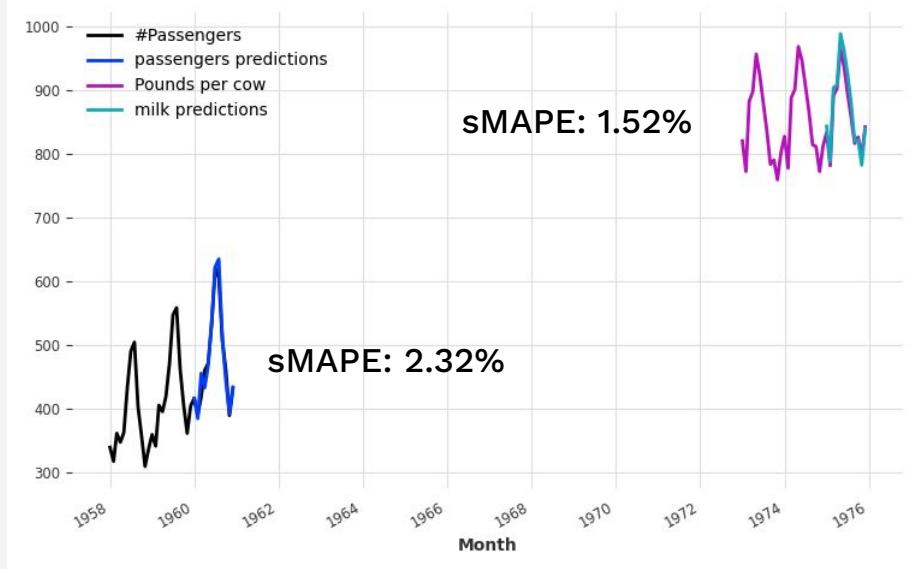
preds = model.predict(
    n=horizon,
    series=train,
    future_covariates=covariates,
)
```



Or just let the model generate these time axis covariates

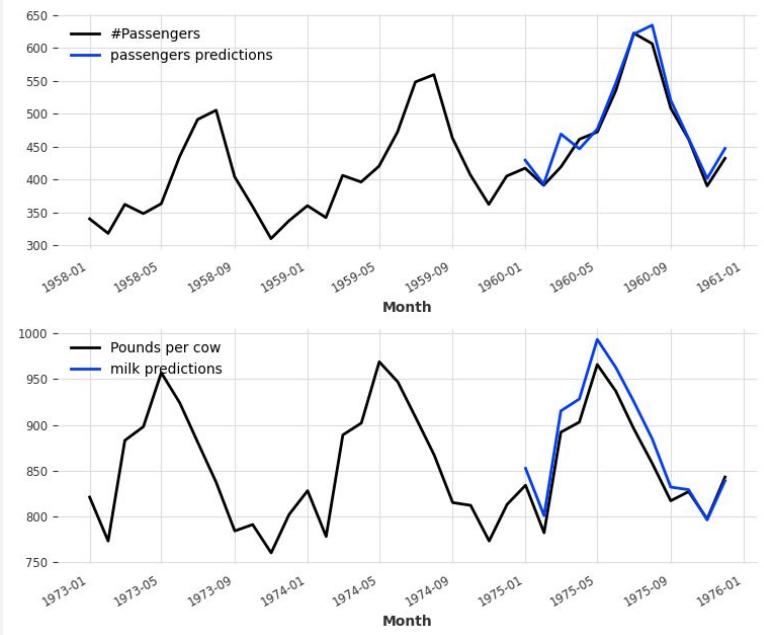
```
# let the model generate the covariates
model = TiDEModel(
    **tide_params,
    use_static_covariates=True,
    add_encoders={"cyclic": {"future": ["month"]}})
)

model.fit(train)
preds = model.predict(n=horizon, series=train)
```



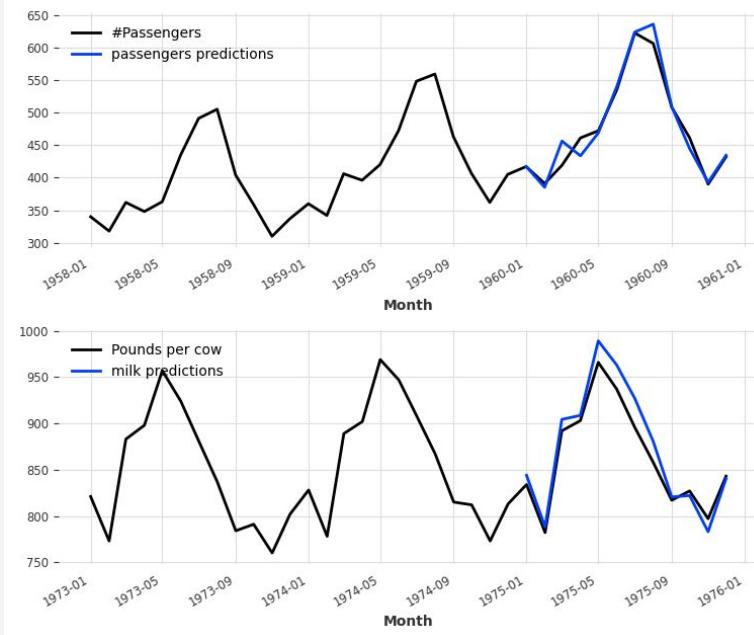
Which one is better?

TiDE - target only



mean sMAPE: 2.46%

TiDE - with covariates



meansMAPE: 1.92% 

TimeSeries

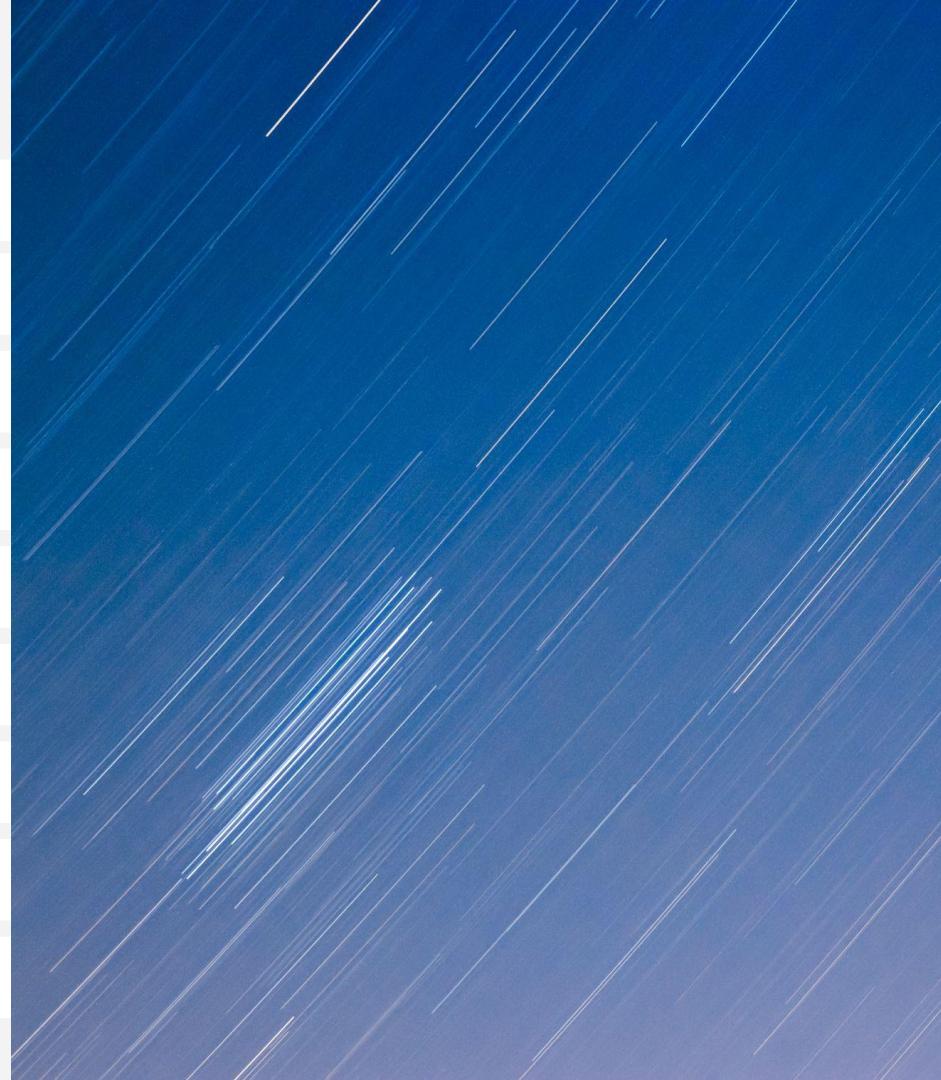
Processing

Forecasting

Evaluating

Agenda

- 1 Intro to Darts
- 2 TimeSeries Class
- 3 Forecasting models
- 4 Simple forecasting example
- 5 Historical forecasting and backtesting
- 6 Training / predicting with multiple series
- 7 Covariates
- 8 Probabilistic Forecasting**
- 9 Summary and what's coming next

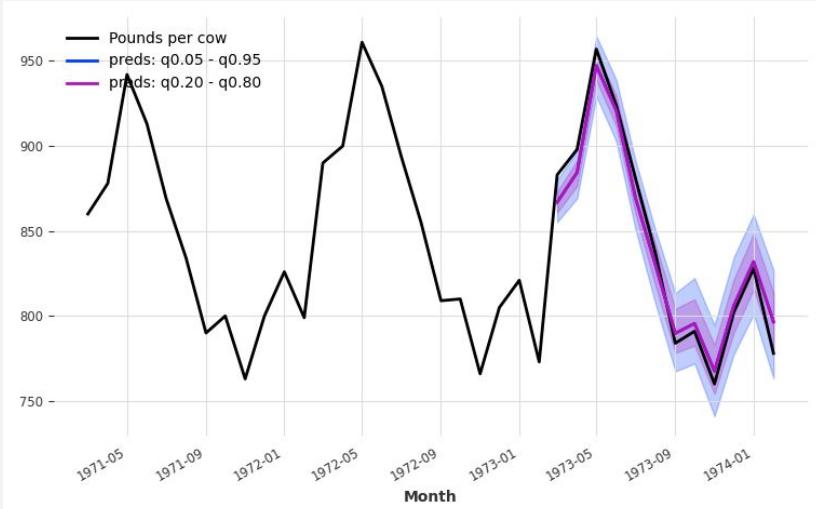


Probabilistic Forecasting with a Local Forecasting Model

```
model = ExponentialSmoothing()
model.fit(train)

# make probabilistic forecasts with num_samples ≥ 1
preds = model.predict(
    n=12,
    num_samples=500,
)

# plotting offers quantile ranges
preds.plot(label="preds: q0.05 - q0.95")
preds.plot(
    label="preds: q0.20 - q0.80",
    low_quantile=0.2,
    high_quantile=0.8
)
```



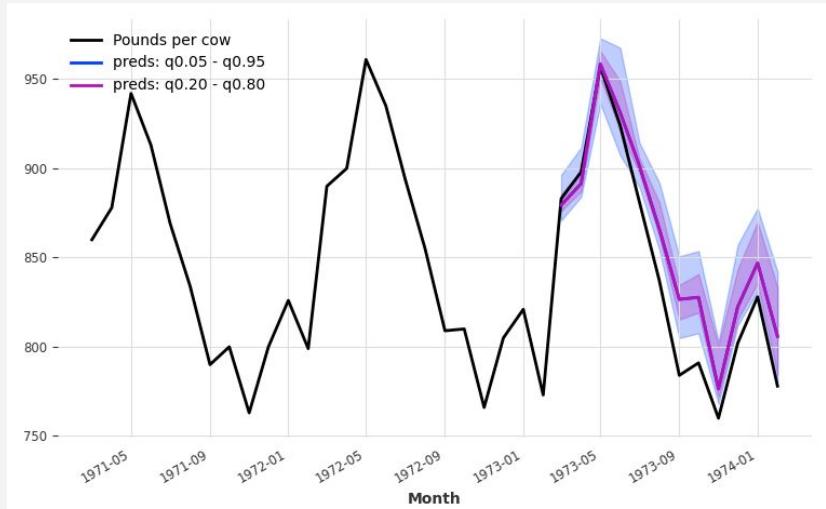
Probabilistic Forecasting with a Global Forecasting Model

```
from darts.utils.likelihood_models import QuantileRegression

# make global models probabilistic by adding a `likelihood`
model = TiDEModel(..., likelihood=QuantileRegression())
model.fit(train)

# make probabilistic forecasts with num_samples ≥ 1
preds = model.predict(
    n=12,
    num_samples=500,
)

# plotting offers quantile ranges
preds.plot(label="preds: q0.05 - q0.95")
preds.plot(
    label="preds: q0.20 - q0.80",
    low_quantile=0.2,
    high_quantile=0.8
)
```



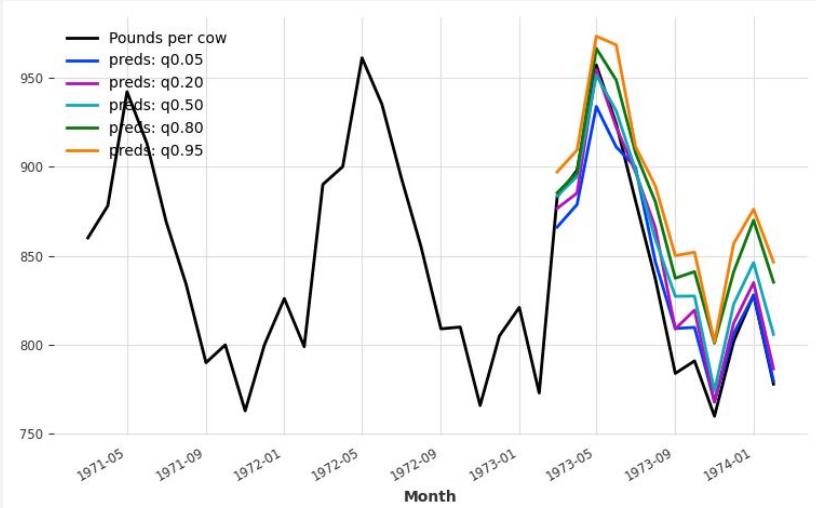
Likelihood Parameter Prediction with a Global Forecasting Model

```
from darts.utils.likelihood_models import QuantileRegression

# make global models probabilistic by adding a `likelihood` parameter
model = TiDEModel(..., likelihood=QuantileRegression())
model.fit(train)

# make probabilistic forecasts with num_samples >= 1
preds = model.predict(
    n=12,
    predict_likelihood_parameters=True,
)

# parameters are stored as columns → plot them directly
preds.plot()
```



Darts' Likelihood Models

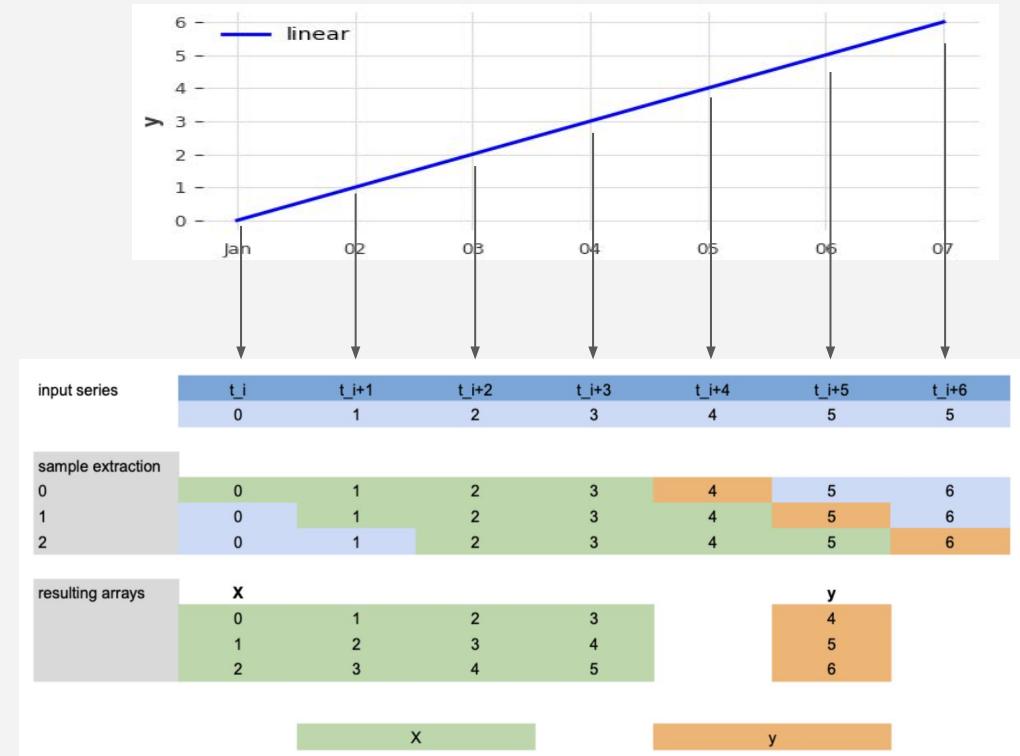
```
from darts.utils.likelihood_models import (
    BernoulliLikelihood,
    BetaLikelihood,
    CauchyLikelihood,
    ContinuousBernoulliLikelihood,
    DirichletLikelihood,
    ExponentialLikelihood,
    GammaLikelihood,
    GaussianLikelihood,
    GeometricLikelihood,
    HalfNormalLikelihood,
    LaplaceLikelihood,
    LogNormalLikelihood,
    NegativeBinomialLikelihood,
    PoissonLikelihood,
    QuantileRegression,
    WeibullLikelihood,
)
```

- + Priors on distributions' parameters



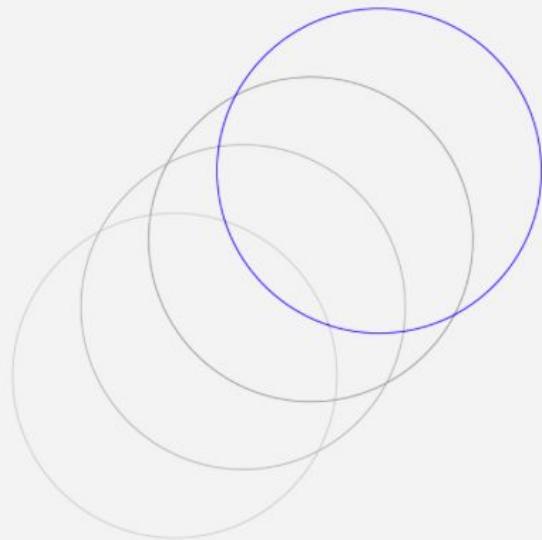
Bonus - Regression Models and Lagged Features

```
from darts import RegressionModel  
  
# create time series  
model = RegressionModel(lags=4, output_chunk_length=1)
```



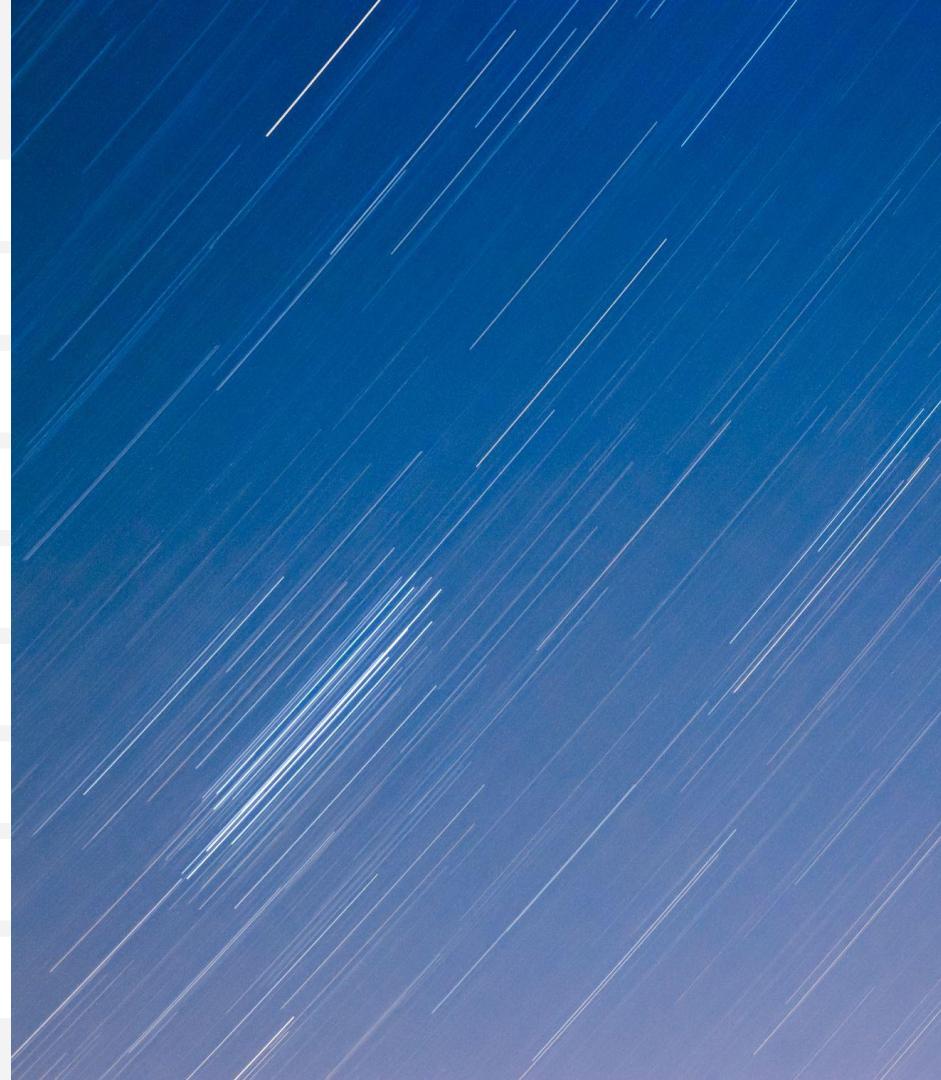
What else exists?

- Anomaly detection
- Explainability (shap for regression models, dedicated explainers for selected neural networks)
- Hierarchical reconciliation
- Filtering
- Dynamic Time Warping
- PyTorch Lightning (multi GPU support, TPU support, callbacks, checkpointing, monitoring, optimizer selection, learning rate scheduling, custom loss functions, ...)
- ...



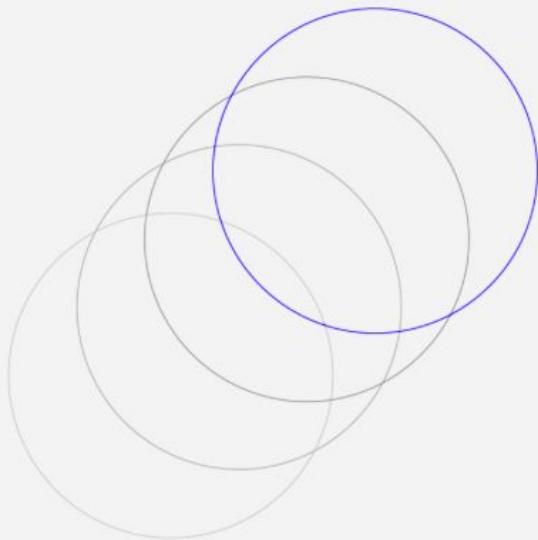
Agenda

- 1 Intro to Darts
- 2 TimeSeries Class
- 3 Forecasting models
- 4 Simple forecasting example
- 5 Historical forecasting and backtesting
- 6 Training / predicting with multiple series
- 7 Covariates
- 8 Probabilistic Forecasting
- 9 Summary and what's coming next**



Take home messages

- Darts is a tool that simplifies time series manipulation and forecasting
- In just a few lines of code, we can use and compare different models
- The unified API allows to use neural network based models in the same way as simpler models
- External data can help to significantly improve forecasts
- There is more than one way of how to forecast (deterministic, probabilistic, global learning on multiple series (the future?))

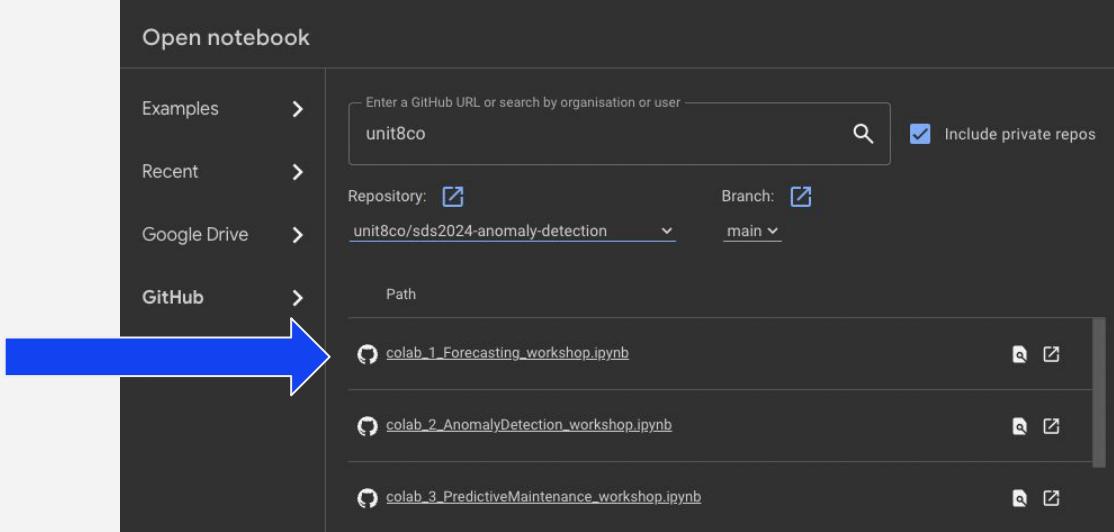


What's planned in the near future?

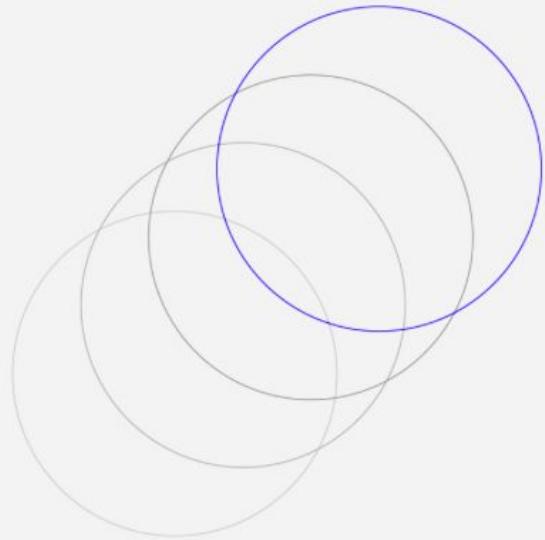
- Conformal Prediction: distribution-free probabilistic forecasting
- Time series Classification

Time To Work

- Open: <https://colab.google/>
- Click on “Open Colab”
- On the left sidebar click GitHub
 - Enter GitHub URL: “unit8co”
 - Select repository: “unit8co/sds2024-anomaly-detection”
 - Select notebook: “colab_1_Forecasting_workshop.ipynb”



Darts



thank you!

